

Assignment 2 - Wine Quality Prediction

1. Data Exploration

1.1 Loading Data

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
from scipy import stats

wine = pd.read_csv("winequality-white.csv", delimiter=';')
```

1.2 Exploratory Data Analysis (EDA)

```
In [2]: wine.head()
```

Out[2]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide	density	pH	sulphates	alcohol
0	7.0	0.27	0.36	20.7	0.045	45.0	170.0	1.0010	3.00	0.45	8.8
1	6.3	0.30	0.34	1.6	0.049	14.0	132.0	0.9940	3.30	0.49	9.5
2	8.1	0.28	0.40	6.9	0.050	30.0	97.0	0.9951	3.26	0.44	10.1
3	7.2	0.23	0.32	8.5	0.058	47.0	186.0	0.9956	3.19	0.40	9.9
4	7.2	0.23	0.32	8.5	0.058	47.0	186.0	0.9956	3.19	0.40	9.9

This dataset contains 11 features and 1 target:

- Features: fixed acidity, volatile acidity, citric acid, residual sugar, chlorides, free sulfur dioxide, total sulfur dioxide, density, pH, sulphates, alcohol
- Target: quality

```
In [3]: wine.info()
```


```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 4898 entries, 0 to 4897
Data columns (total 12 columns):
#   Column                Non-Null Count  Dtype
---  -
0   fixed acidity          4898 non-null   float64
1   volatile acidity       4898 non-null   float64
2   citric acid            4898 non-null   float64
3   residual sugar         4898 non-null   float64
4   chlorides              4898 non-null   float64
5   free sulfur dioxide    4898 non-null   float64
6   total sulfur dioxide   4898 non-null   float64
7   density                4898 non-null   float64
8   pH                    4898 non-null   float64
9   sulphates              4898 non-null   float64
10  alcohol                4898 non-null   float64
11  quality                4898 non-null   int64
dtypes: float64(11), int64(1)
memory usage: 459.3 KB
```

There are 4898 entries and all are numerical (long float and long integer). Plus there is no null value in all of the columns.

```
In [4]: wine.describe()
```

Out[4]:

	fixed acidity	volatile acidity	citric acid	residual sugar	chlorides	free sulfur dioxide	total sulfur dioxide
count	4898.000000	4898.000000	4898.000000	4898.000000	4898.000000	4898.000000	4898.000000
mean	6.854788	0.278241	0.334192	6.391415	0.045772	35.308085	138.36065
std	0.843868	0.100795	0.121020	5.072058	0.021848	17.007137	42.49806
min	3.800000	0.080000	0.000000	0.600000	0.009000	2.000000	9.00000
25%	6.300000	0.210000	0.270000	1.700000	0.036000	23.000000	108.00000
50%	6.800000	0.260000	0.320000	5.200000	0.043000	34.000000	134.00000
75%	7.300000	0.320000	0.390000	9.900000	0.050000	46.000000	167.00000
max	14.200000	1.100000	1.660000	65.800000	0.346000	289.000000	440.00000




It seems like there are outliers in the residual sugar feature. We can see that 25% of the values are under 1.70 and 75% of them are under 9.90. Taking into account the mean at 6.39, the minimum value of 0.60 and maximum value of 65.80 seems like an outlier.

Using boxplot to visualize the outlier in residual sugar.

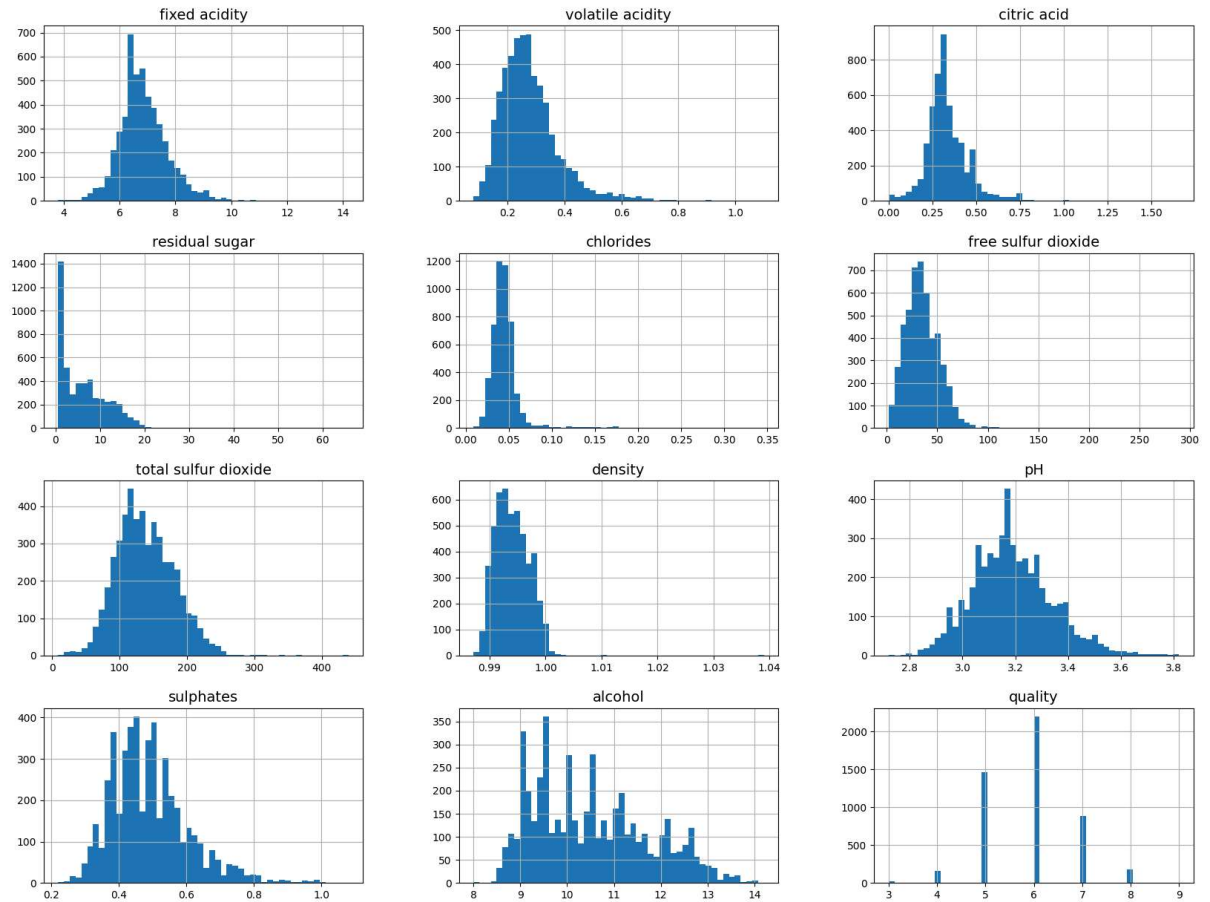
```
In [5]: sns.boxplot(x=wine['residual sugar'])
```

```
Out[5]: <AxesSubplot: xlabel='residual sugar'>
```



It seems like residual sugar > 20 is considered outlier.

```
In [6]: plt.rc('font', size=14)
plt.rc('axes', labels=14, titles=14)
plt.rc('legend', font=14)
plt.rc('xtick', labels=10)
plt.rc('ytick', labels=10)
wine.hist(bins=50, figsize=(20,15))
plt.show()
```



Observations:

- most features are normally distributed except for residual sugar and alcohol
- a lot of features are skewed to the right (fixed acidity, volatile acidity, citric acid, residual sugar, chlorides, free sulfur dioxide, total sulfur dioxide, density)
- from quality, we can see that we have data imbalance problem where we might not have enough samples of high quality wine.

Checking class imbalance.


```
In [7]: wine['quality'].value_counts()
```

```
Out[7]: quality
6      2198
5      1457
7       880
8       175
4       163
3        20
9         5
Name: count, dtype: int64
```

1.3 Visualize Correlation between Features and Target

```
In [8]: corr = wine.corr()

plt.figure(figsize=(12, 12))
sns.heatmap(corr, annot=True, fmt='.2f', cmap='coolwarm', square=True, cbar=False)
b, t = plt.ylim()
plt.ylim(b+0.5, t-0.5)
plt.title("Feature Correlation Heatmap")
plt.show()
```



We can see that alcohol has the highest correlation with the wine quality followed by density, chlorides, and volatile acidity.

2. Data Preprocessing

2.1 Handling Outliers

Using z-score to find and remove outliers with threshold set to 3.

```
In [9]: z = np.abs(stats.zscore(wine))
wine = wine[(z < 3).all(axis=1)]
print(wine.shape)

wine['quality'].value_counts()

(4487, 12)
```

```
Out[9]: quality
6      2038
5      1309
7       855
8       161
4       124
Name: count, dtype: int64
```

After removing outliers, there are 4487 entries left. This means there are about 411 outliers (about 8.4%).

2.2 Split Features and Target

```
In [10]: X = wine.drop('quality', axis=1)
print(X.shape)

y = wine['quality']
print(y.shape)

(4487, 11)
(4487,)
```

2.3 Perform Feature Scaling

```
In [11]: from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X = scaler.fit_transform(X)
```

2.4 Split Training and Test Data

```
In [12]: from sklearn.model_selection import train_test_split

X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2, random_state=42, stratify=y)
print(X_train.shape)
print(X_test.shape)
print(y_train.shape)
print(y_test.shape)

(3589, 11)
(898, 11)
(3589,)
(898,)
```

3. Binary Classification

3.1 Categorizing Wine Quality

For binary classification task:

- high quality ≥ 7 quality rating
- low quality < 7 quality rating

```
In [13]: y_train_hq = (y_train >= 7)
y_test_hq = (y_test >= 7)
```

3.2 Model Training

3.2.1 Logistic Regression


```
In [14]: from sklearn.linear_model import LogisticRegression

logreg = LogisticRegression()
logreg.fit(X_train, y_train_hq)
```

```
Out[14]: LogisticRegression
LogisticRegression()
```

3.2.2 Support Vector Machine (SVM)

```
In [15]: from sklearn.svm import SVC

svm = SVC(random_state=42)
svm.fit(X_train, y_train_hq)
```

```
Out[15]: SVC
SVC(random_state=42)
```

3.2.3 K-Nearest Neighbors (KNN)

Before training, lets see how the accuracy changed with number of K.

```

In [16]: from sklearn.neighbors import KNeighborsClassifier
from sklearn.model_selection import cross_val_score


# Range of k from 1 to 26
k_range = range(1, 26)

# List of cv scores
cv_scores = []

# Perform 10-fold cross validation
for k in k_range:
    knn = KNeighborsClassifier(n_neighbors=k)
    scores = cross_val_score(knn, X, y, cv=10, scoring='accuracy')
    cv_scores.append(scores.mean())

# Plot accuracy vs k
plt.plot(k_range, cv_scores)
plt.xlabel('k')
plt.ylabel('Accuracy')
plt.title('Accuracy vs k')
plt.show()

```



We can see that highest accuracy is achieved when k equals to 19 or 23. Train the model with k = 23.

```
In [17]: knn = KNeighborsClassifier(n_neighbors=23)
knn.fit(X_train, y_train_hq)
```

```
Out[17]: KNeighborsClassifier
KNeighborsClassifier(n_neighbors=23)
```

3.2.4 Decision Tree

```
In [18]: from sklearn.tree import DecisionTreeClassifier

dt = DecisionTreeClassifier(random_state=42)
dt.fit(X_train, y_train_hq)
```

```
Out[18]: DecisionTreeClassifier
DecisionTreeClassifier(random_state=42)
```

3.2.5 Random Forest

```
In [19]: from sklearn.ensemble import RandomForestClassifier

rf = RandomForestClassifier(random_state=42)
rf.fit(X_train, y_train_hq)
```

```
Out[19]: RandomForestClassifier
RandomForestClassifier(random_state=42)
```

3.3 Model Evaluation

In model evaluation, the trained model will be cross-validated using 10-fold and tested using the test set.

```
In [20]: # Cross-validation using 10-fold with accuracy scoring
y_cv_logreg = cross_val_score(logreg, X_train, y_train_hq, cv=10, scoring='accuracy')
y_cv_svm = cross_val_score(svm, X_train, y_train_hq, cv=10, scoring='accuracy')
y_cv_knn = cross_val_score(knn, X_train, y_train_hq, cv=10, scoring='accuracy')
y_cv_dt = cross_val_score(dt, X_train, y_train_hq, cv=10, scoring='accuracy')
y_cv_rf = cross_val_score(rf, X_train, y_train_hq, cv=10, scoring='accuracy')

# Predict using the test set
y_pred_logreg = logreg.predict(X_test)
y_pred_svm = svm.predict(X_test)
y_pred_knn = knn.predict(X_test)
y_pred_dt = dt.predict(X_test)
y_pred_rf = rf.predict(X_test)
```

3.3.1 Evaluation of Logistic Regression


```
In [21]: from sklearn.metrics import classification_report, ConfusionMatrixDisplay

print(classification_report(y_test_hq, y_pred_logreg, digits=3, zero_division=
1))
print('Cross-validation accuracy score:', y_cv_logreg.mean())

plt.rc('font', size=14)
ConfusionMatrixDisplay.from_predictions(y_test_hq, y_pred_logreg)
plt.show()
```

	precision	recall	f1-score	support
False	0.824	0.934	0.875	695
True	0.582	0.315	0.409	203
accuracy			0.794	898
macro avg	0.703	0.625	0.642	898
weighted avg	0.769	0.794	0.770	898

Cross-validation accuracy score: 0.8013312895846626




3.3.2 Evaluation of SVM

```
In [22]: print(classification_report(y_test_hq, y_pred_svm, digits=3, zero_division=1))
print('Cross-validation accuracy score:', y_cv_svm.mean())

plt.rc('font', size=14)
ConfusionMatrixDisplay.from_predictions(y_test_hq, y_pred_svm)
plt.show()
```

	precision	recall	f1-score	support
False	0.840	0.944	0.889	695
True	0.667	0.384	0.487	203
accuracy			0.817	898
macro avg	0.753	0.664	0.688	898
weighted avg	0.801	0.817	0.798	898

Cross-validation accuracy score: 0.8183338261153732




3.3.3 Evaluation of KNN

```
In [23]: print(classification_report(y_test_hq, y_pred_knn, digits=3, zero_division=1))
print('Cross-validation accuracy score:', y_cv_knn.mean())

plt.rc('font', size=14)
ConfusionMatrixDisplay.from_predictions(y_test_hq, y_pred_knn)
plt.show()
```

	precision	recall	f1-score	support
False	0.854	0.917	0.884	695
True	0.618	0.463	0.530	203
accuracy			0.814	898
macro avg	0.736	0.690	0.707	898
weighted avg	0.801	0.814	0.804	898

Cross-validation accuracy score: 0.8144317704361898




3.3.4 Evaluation of Decision Tree

```
In [24]: print(classification_report(y_test_hq, y_pred_dt, digits=3, zero_division=1))
print('Cross-validation accuracy score:', y_cv_dt.mean())

plt.rc('font', size=14)
ConfusionMatrixDisplay.from_predictions(y_test_hq, y_pred_dt)
plt.show()
```

	precision	recall	f1-score	support
False	0.893	0.845	0.868	695
True	0.552	0.655	0.599	203
accuracy			0.802	898
macro avg	0.723	0.750	0.734	898
weighted avg	0.816	0.802	0.807	898

Cross-validation accuracy score: 0.8339342680630553




3.3.5 Evaluation of Random Forest


```
In [25]: print(classification_report(y_test_hq, y_pred_rf, digits=3, zero_division=1))
print('Cross-validation accuracy score:', y_cv_rf.mean())

plt.rc('font', size=14)
ConfusionMatrixDisplay.from_predictions(y_test_hq, y_pred_rf)
plt.show()
```

	precision	recall	f1-score	support
False	0.894	0.947	0.920	695
True	0.772	0.616	0.685	203
accuracy			0.872	898
macro avg	0.833	0.781	0.802	898
weighted avg	0.866	0.872	0.867	898

Cross-validation accuracy score: 0.873226373694776



Ranking:

- CV Accuracy : RF, DT, SVM, KNN, LOGREG
- Precision : RF, DT, KNN, SVM, LOGREG
- Recall : RF, KNN, DT, SVM, LOGREG
- F1-Score : RF, DT, KNN, SVM, LOGREG

Observations:

- Random Forest is the best at binary classification for this white wine quality prediction.
- However, Decision Tree is better at recall high quality wine.

3.4 ROC Curve Analysis

```
In [26]: from sklearn.metrics import roc_curve, auc


models = [logreg, svm, knn, dt, rf]
model_names = ['Logistic Regression', 'SVM', 'KNN', 'Decision Tree', 'Random Forest']

plt.figure(figsize=(8, 6))

# Plot ROC curve for each model
for model, name in zip(models, model_names):
    y_pred = model.predict(X_test)
    fpr, tpr, thresholds = roc_curve(y_test_hq, y_pred)
    roc_auc = auc(fpr, tpr)
    plt.plot(fpr, tpr, label='%s (AUC = %0.2f)' % (name, roc_auc))

# Plot the random guess line
plt.plot([0, 1], [0, 1], 'k--')


plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('ROC Curve for Binary Classification')
plt.legend(loc="lower right")
plt.show()
```



From the ROC curve, we can observe that Random Forest have the highest ROC AUC which is 0.78. Although this value is quite low (ROC AUC = 1 being the perfect), it is the best we can achieved among the 5 chosen algorithms with imbalanced dataset.

3.5 Feature Importance for Binary Classification

```
In [27]: feature_importances_ = rf.feature_importances_  
  
# Plot the feature importance for Random Forest  
plt.figure(figsize=(10, 6))  
plt.barh(wine.columns[:-1], feature_importances_)  
plt.xlabel('Feature Importance')  
plt.ylabel('Feature')  
plt.title('Feature Importance for Random Forest')  
plt.show()
```



For binary classification, alcohol and density are the most important features to distinguish low quality wine and high quality wine (binary). This might not be the case for multiclass classification.

4. Multiclass Classification

4.1 Model Training

Logistic Regression and Support Vector Machine, by default, is limited to binary classification problems. However, we can utilize extensions like One-vs-One (OvO) and One-vs-Rest (OvR) to be used for multiclass classification problems. Other algorithms such as K-Nearest Neighbors, Decision Tree, and Random Forest inherently support multiclass classification.

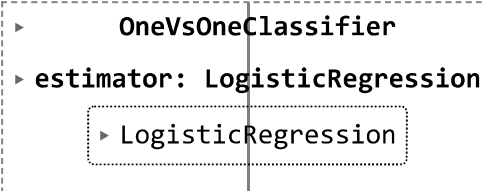
4.1.1 Logistic Regression

One-vs-One approach.

```
In [28]: from sklearn.multiclass import OneVsOneClassifier

# One-vs-One
ovo_logreg = OneVsOneClassifier(LogisticRegression())
ovo_logreg.fit(X_train, y_train)
```

```
Out[28]:
```



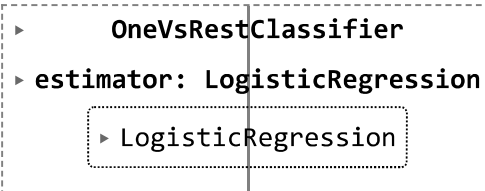
```
  ▸ OneVsOneClassifier
    ▸ estimator: LogisticRegression
      ▸ LogisticRegression
```

One-vs-Rest approach.

```
In [29]: from sklearn.multiclass import OneVsRestClassifier

# One-vs-Rest
ovr_logreg = OneVsRestClassifier(LogisticRegression())
ovr_logreg.fit(X_train, y_train)
```

```
Out[29]:
```



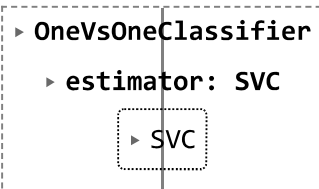
```
  ▸ OneVsRestClassifier
    ▸ estimator: LogisticRegression
      ▸ LogisticRegression
```

4.1.2 Support Vector Machines (SVM)

One-vs-One approach.

```
In [30]: ovo_svm = OneVsOneClassifier(SVC(random_state=42))
ovo_svm.fit(X_train, y_train)
```

```
Out[30]:
```



```
  ▸ OneVsOneClassifier
    ▸ estimator: SVC
      ▸ SVC
```

One-vs-Rest approach.

```
In [31]: ovr_svm = OneVsRestClassifier(SVC(random_state=42))  
ovr_svm.fit(X_train, y_train)
```

```
Out[31]: ▸ OneVsRestClassifier  
          ▸ estimator: SVC  
            ▸ SVC
```

4.1.3 K-Nearest Neighbors (KNN)

```
In [32]: knn = KNeighborsClassifier(n_neighbors=23)  
knn.fit(X_train, y_train)
```

```
Out[32]: ▾ KNeighborsClassifier  
          KNeighborsClassifier(n_neighbors=23)
```

4.1.4 Decision Tree

```
In [33]: dt = DecisionTreeClassifier(random_state=42)  
dt.fit(X_train, y_train)
```

```
Out[33]: ▾ DecisionTreeClassifier  
          DecisionTreeClassifier(random_state=42)
```

4.1.5 Random Forest

```
In [34]: rf = RandomForestClassifier(random_state=42)  
rf.fit(X_train, y_train)
```

```
Out[34]: ▾ RandomForestClassifier  
          RandomForestClassifier(random_state=42)
```

4.2 Model Evaluation

For multiclass classification, it is recommended to find AUC score using cross validation method by passing `roc_auc_ovr` parameter. Evaluating the model with this method is advantageous when there is a high class imbalance. Also it does not require to set a classification threshold.

```
In [35]: # Cross-validation using 10-fold with accuracy scoring
y_cv_ovo_logreg = cross_val_score(ovo_logreg, X_train, y_train, cv=10, scoring
='accuracy')
y_cv_ovr_logreg = cross_val_score(ovr_logreg, X_train, y_train, cv=10, scoring
='accuracy')
y_cv_ovo_svm = cross_val_score(ovo_svm, X_train, y_train, cv=10, scoring='accu
racy')
y_cv_ovr_svm = cross_val_score(ovr_svm, X_train, y_train, cv=10, scoring='accu
racy')
y_cv_knn = cross_val_score(knn, X_train, y_train, cv=10, scoring='accuracy')
y_cv_dt = cross_val_score(dt, X_train, y_train, cv=10, scoring='accuracy')
y_cv_rf = cross_val_score(rf, X_train, y_train, cv=10, scoring='accuracy')

# Predict using the test set
y_pred_ovo_logreg = ovo_logreg.predict(X_test)
y_pred_ovr_logreg = ovr_logreg.predict(X_test)
y_pred_ovo_svm = ovo_svm.predict(X_test)
y_pred_ovr_svm = ovr_svm.predict(X_test)
y_pred_knn = knn.predict(X_test)
y_pred_dt = dt.predict(X_test)
y_pred_rf = rf.predict(X_test)
```

4.2.1 Evaluation of Logistic Regression


One-vs-One.


```
In [36]: print(classification_report(y_test, y_pred_ovo_logreg, digits=3, zero_division=1))
print('Cross-validation accuracy score:', y_cv_ovo_logreg.mean())

plt.rc('font', size=14)
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_ovo_logreg)
plt.show()
```

	precision	recall	f1-score	support
4	1.000	0.000	0.000	25
5	0.578	0.481	0.525	262
6	0.527	0.752	0.620	408
7	0.500	0.287	0.364	171
8	1.000	0.000	0.000	32
accuracy			0.537	898
macro avg	0.721	0.304	0.302	898
weighted avg	0.567	0.537	0.504	898

Cross-validation accuracy score: 0.5391473833273681




One-vs-Rest.

```
In [37]: print(classification_report(y_test, y_pred_ovr_logreg, digits=3, zero_division=1))
print('Cross-validation accuracy score:', y_cv_ovr_logreg.mean())

plt.rc('font', size=14)
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_ovr_logreg)
plt.show()
```

	precision	recall	f1-score	support
4	1.000	0.000	0.000	25
5	0.583	0.481	0.527	262
6	0.505	0.762	0.607	408
7	0.515	0.199	0.287	171
8	1.000	0.000	0.000	32
accuracy			0.524	898
macro avg	0.721	0.288	0.284	898
weighted avg	0.561	0.524	0.484	898

Cross-validation accuracy score: 0.5469476043012091



Although the OvR method achieved higher cross-validation accuracy, the OvO achieved better performance in all aspects (accuracy, precision, recall, f1-score) during testing.

4.2.2 Evaluation of SVM


One-vs-One.

```
In [38]: print(classification_report(y_test, y_pred_ovo_svm, digits=3, zero_division=
1))
print('Cross-validation accuracy score:', y_cv_ovo_svm.mean())

plt.rc('font', size=14)
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_ovo_svm)
plt.show()
```

	precision	recall	f1-score	support
4	1.000	0.000	0.000	25
5	0.610	0.561	0.584	262
6	0.548	0.760	0.637	408
7	0.571	0.304	0.397	171
8	1.000	0.000	0.000	32
accuracy			0.567	898
macro avg	0.746	0.325	0.324	898
weighted avg	0.599	0.567	0.535	898

Cross-validation accuracy score: 0.579545136241266




One-vs-Rest.

```
In [39]: print(classification_report(y_test, y_pred_ovr_svm, digits=3, zero_division=
1))
print('Cross-validation accuracy score:', y_cv_ovr_svm.mean())

plt.rc('font', size=14)
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_ovr_svm)
plt.show()
```

	precision	recall	f1-score	support
4	1.000	0.000	0.000	25
5	0.610	0.580	0.595	262
6	0.554	0.728	0.629	408
7	0.527	0.339	0.413	171
8	1.000	0.094	0.171	32
accuracy			0.568	898
macro avg	0.738	0.348	0.362	898
weighted avg	0.594	0.568	0.544	898

Cross-validation accuracy score: 0.5739756617544078



SVM however achieved better performance using the OvR method. Plus, it is the first time a model is able to predict samples with quality class 8.


4.2.3 Evaluation of KNN

```
In [40]: print(classification_report(y_test, y_pred_knn, digits=3, zero_division=1))
print('Cross-validation accuracy score:', y_cv_knn.mean())

plt.rc('font', size=14)
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_knn)
plt.show()
```

	precision	recall	f1-score	support
4	1.000	0.000	0.000	25
5	0.588	0.584	0.586	262
6	0.557	0.679	0.612	408
7	0.479	0.392	0.431	171
8	0.000	0.000	0.000	32
accuracy			0.553	898
macro avg	0.525	0.331	0.326	898
weighted avg	0.544	0.553	0.531	898

Cross-validation accuracy score: 0.5550240425763682




4.2.4 Evaluation of Decision Tree

```
In [41]: print(classification_report(y_test, y_pred_dt, digits=3, zero_division=1))
print('Cross-validation accuracy score:', y_cv_dt.mean())

plt.rc('font', size=14)
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_dt)
plt.show()
```

	precision	recall	f1-score	support
4	0.455	0.400	0.426	25
5	0.629	0.634	0.631	262
6	0.662	0.637	0.649	408
7	0.560	0.596	0.578	171
8	0.324	0.375	0.348	32
accuracy			0.612	898
macro avg	0.526	0.528	0.526	898
weighted avg	0.615	0.612	0.613	898

Cross-validation accuracy score: 0.6154946234885856




4.2.5 Evaluation of Random Forest

```
In [42]: print(classification_report(y_test, y_pred_rf, digits=3, zero_division=1))
print('Cross-validation accuracy score:', y_cv_rf.mean())

plt.rc('font', size=14)
ConfusionMatrixDisplay.from_predictions(y_test, y_pred_rf)
plt.show()
```

	precision	recall	f1-score	support
4	1.000	0.320	0.485	25
5	0.733	0.649	0.688	262
6	0.640	0.767	0.698	408
7	0.647	0.591	0.618	171
8	0.923	0.375	0.533	32
accuracy			0.673	898
macro avg	0.789	0.540	0.604	898
weighted avg	0.689	0.673	0.668	898


Cross-validation accuracy score: 0.6904452156051104



Random forest is the best algorithm for the multiclass classification of wine quality. It achieved the highest cross-validation accuracy, testing accuracy, precision, recall, and f1-score. However, decision tree is better at recall for minority classes. This is the exact case as in binary classification.

4.3 Feature Importance for Multiclass Classification

```
In [43]: feature_importances_ = rf.feature_importances_  
  
# Plot the feature importance for Random Forest  
plt.figure(figsize=(10, 6))  
plt.barh(wine.columns[:-1], feature_importances_)  
plt.xlabel('Feature Importance')  
plt.ylabel('Feature')  
plt.title('Feature Importance for Random Forest')  
plt.show()
```



For multiclass classification, alcohol and density are still the most important features as in binary classification. However, we can also see that volatile acidity has high importance during training.