

Neural Network: Regression using MLP in Keras

1. Data Loading and Preparation

Import required dependencies

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import tensorflow as tf

import warnings
warnings.filterwarnings('ignore')
```

Load California Housing dataset from TensorFlow API. Use 60:20:20 ratio for the segmentation of training, validation, and testing datasets

```
(X_train_full, y_train_full), (X_test, y_test) =
tf.keras.datasets.california_housing.load_data(
    path='california_housing.npz',
    test_split=0.2,
    seed=113
)

from sklearn.model_selection import train_test_split

# 0.25 * 0.8 = 0.2
X_train, X_valid, y_train, y_valid = train_test_split(X_train_full,
y_train_full, test_size=0.25)

print("Training shape:", X_train.shape, y_train.shape)
print("Validation shape:", X_valid.shape, y_valid.shape)
print("Testing shape:", X_test.shape, y_test.shape)

Training shape: (12384, 8) (12384,)
Validation shape: (4128, 8) (4128,)
Testing shape: (4128, 8) (4128,)
```

2. Data Preprocessing

From previous assignment, the median_housing_value (our target) contains outlier at value more than 500,000. It is observed that the dataset clipped any houses with price more than 500,000 (threshold). Thus, if we plot the histogram, we will see a huge surge when the median_housing_value is more than 500,000.

Let's remove this outliers.

```
X_train = X_train[y_train < 500000]
y_train = y_train[y_train < 500000]
X_test = X_test[y_test < 500000]
y_test = y_test[y_test < 500000]
X_valid = X_valid[y_valid < 500000]
y_valid = y_valid[y_valid < 500000]
```

Standardize the features using StandardScaler.

```
from sklearn.preprocessing import StandardScaler

scaler = StandardScaler()
X_train = scaler.fit_transform(X_train)
X_valid = scaler.transform(X_valid)
X_test = scaler.transform(X_test)
```

Scale the target variable to ensure it falls within a specific range, which can help the model to learn more effectively.

```
y_train = scaler.fit_transform(y_train.reshape(-1, 1)).flatten()
y_valid = scaler.transform(y_valid.reshape(-1, 1)).flatten()
y_test = scaler.transform(y_test.reshape(-1, 1)).flatten()
```

3. Build Sequential Regression MLP Model

Model Architecture:

- Input Layer: Takes input of shape equal to the number of features (8).
- First Hidden Layer: Fully connected layer with 128 neurons and ReLU activation.
- Dropout Layer: Applies dropout to 20% of the neurons to prevent overfitting.
- Second Hidden Layer: Fully connected layer with 64 neurons and ReLU activation.
- Output Layer: Fully connected layer with 1 neuron, providing the final regression output.

```
model = tf.keras.models.Sequential([
    tf.keras.layers.Dense(128, activation='relu',
input_shape=(X_train.shape[1],)),
    tf.keras.layers.Dropout(0.2), # Dropout for regularization
    tf.keras.layers.Dense(64, activation='relu'),
    tf.keras.layers.Dense(1)
])
```

3. Compile Model

- Use MSE as loss function with Adam optimizer
- Evaluate model using MAE
- Additional: Learning rate of the adam optimizer is adjusted

```
model.compile(optimizer=tf.keras.optimizers.Adam(learning_rate=0.001),  
loss='mean_squared_error', metrics=['mean_absolute_error'])
```

4. Model Training

```
history = model.fit(X_train, y_train, epochs=100, batch_size=32,  
validation_data=(X_valid, y_valid))
```

Epoch 1/100

369/369 _____ 1s 1ms/step - loss: 0.5337 -
mean_absolute_error: 0.5497 - val_loss: 0.3317 -
val_mean_absolute_error: 0.4237

Epoch 2/100

369/369 _____ 0s 960us/step - loss: 0.3532 -
mean_absolute_error: 0.4398 - val_loss: 0.3108 -
val_mean_absolute_error: 0.3984

Epoch 3/100

369/369 _____ 0s 1ms/step - loss: 0.3239 -
mean_absolute_error: 0.4156 - val_loss: 0.2990 -
val_mean_absolute_error: 0.3904

Epoch 4/100

369/369 _____ 0s 965us/step - loss: 0.3043 -
mean_absolute_error: 0.3979 - val_loss: 0.3011 -
val_mean_absolute_error: 0.4007

Epoch 5/100

369/369 _____ 0s 1ms/step - loss: 0.3091 -
mean_absolute_error: 0.4010 - val_loss: 0.2909 -
val_mean_absolute_error: 0.3732

Epoch 6/100

369/369 _____ 1s 2ms/step - loss: 0.2935 -
mean_absolute_error: 0.3921 - val_loss: 0.2887 -
val_mean_absolute_error: 0.3824

Epoch 7/100

369/369 _____ 0s 981us/step - loss: 0.2795 -
mean_absolute_error: 0.3814 - val_loss: 0.2726 -
val_mean_absolute_error: 0.3670

Epoch 8/100

369/369 _____ 0s 989us/step - loss: 0.2820 -
mean_absolute_error: 0.3823 - val_loss: 0.2712 -
val_mean_absolute_error: 0.3741

Epoch 9/100

369/369 _____ 0s 998us/step - loss: 0.2685 -
mean_absolute_error: 0.3712 - val_loss: 0.2598 -
val_mean_absolute_error: 0.3508

Epoch 10/100

369/369 _____ 0s 982us/step - loss: 0.2684 -
mean_absolute_error: 0.3710 - val_loss: 0.2628 -
val_mean_absolute_error: 0.3606

Epoch 11/100

```
369/369 _____ 0s 1ms/step - loss: 0.2691 -  
mean_absolute_error: 0.3723 - val_loss: 0.2554 -  
val_mean_absolute_error: 0.3506  
Epoch 12/100  
369/369 _____ 0s 1ms/step - loss: 0.2598 -  
mean_absolute_error: 0.3652 - val_loss: 0.2535 -  
val_mean_absolute_error: 0.3489  
Epoch 13/100  
369/369 _____ 1s 2ms/step - loss: 0.2577 -  
mean_absolute_error: 0.3613 - val_loss: 0.2534 -  
val_mean_absolute_error: 0.3446  
Epoch 14/100  
369/369 _____ 1s 1ms/step - loss: 0.2530 -  
mean_absolute_error: 0.3572 - val_loss: 0.2559 -  
val_mean_absolute_error: 0.3563  
Epoch 15/100  
369/369 _____ 0s 1ms/step - loss: 0.2541 -  
mean_absolute_error: 0.3585 - val_loss: 0.2485 -  
val_mean_absolute_error: 0.3429  
Epoch 16/100  
369/369 _____ 1s 1ms/step - loss: 0.2477 -  
mean_absolute_error: 0.3542 - val_loss: 0.2518 -  
val_mean_absolute_error: 0.3534  
Epoch 17/100  
369/369 _____ 0s 1ms/step - loss: 0.2572 -  
mean_absolute_error: 0.3572 - val_loss: 0.2459 -  
val_mean_absolute_error: 0.3400  
Epoch 18/100  
369/369 _____ 0s 1ms/step - loss: 0.2558 -  
mean_absolute_error: 0.3582 - val_loss: 0.2455 -  
val_mean_absolute_error: 0.3409  
Epoch 19/100  
369/369 _____ 0s 981us/step - loss: 0.2491 -  
mean_absolute_error: 0.3513 - val_loss: 0.2453 -  
val_mean_absolute_error: 0.3427  
Epoch 20/100  
369/369 _____ 0s 1ms/step - loss: 0.2496 -  
mean_absolute_error: 0.3535 - val_loss: 0.2438 -  
val_mean_absolute_error: 0.3387  
Epoch 21/100  
369/369 _____ 0s 1ms/step - loss: 0.2405 -  
mean_absolute_error: 0.3472 - val_loss: 0.2394 -  
val_mean_absolute_error: 0.3417  
Epoch 22/100  
369/369 _____ 0s 1ms/step - loss: 0.2366 -  
mean_absolute_error: 0.3442 - val_loss: 0.2474 -  
val_mean_absolute_error: 0.3486  
Epoch 23/100  
369/369 _____ 0s 1ms/step - loss: 0.2284 -
```

```
mean_absolute_error: 0.3443 - val_loss: 0.2433 -  
val_mean_absolute_error: 0.3402  
Epoch 24/100  
369/369 ━━━━━━━━━━━ 0s 1ms/step - loss: 0.2343 -  
mean_absolute_error: 0.3433 - val_loss: 0.2446 -  
val_mean_absolute_error: 0.3465  
Epoch 25/100  
369/369 ━━━━━━━━━━━ 0s 1ms/step - loss: 0.2484 -  
mean_absolute_error: 0.3530 - val_loss: 0.2357 -  
val_mean_absolute_error: 0.3365  
Epoch 26/100  
369/369 ━━━━━━━━━━━ 0s 980us/step - loss: 0.2423 -  
mean_absolute_error: 0.3477 - val_loss: 0.2406 -  
val_mean_absolute_error: 0.3452  
Epoch 27/100  
369/369 ━━━━━━━━━━━ 0s 1ms/step - loss: 0.2292 -  
mean_absolute_error: 0.3357 - val_loss: 0.2383 -  
val_mean_absolute_error: 0.3458  
Epoch 28/100  
369/369 ━━━━━━━━━━━ 0s 982us/step - loss: 0.2406 -  
mean_absolute_error: 0.3450 - val_loss: 0.2342 -  
val_mean_absolute_error: 0.3324  
Epoch 29/100  
369/369 ━━━━━━━━━━━ 0s 1ms/step - loss: 0.2399 -  
mean_absolute_error: 0.3444 - val_loss: 0.2410 -  
val_mean_absolute_error: 0.3357  
Epoch 30/100  
369/369 ━━━━━━━━━━━ 0s 1ms/step - loss: 0.2412 -  
mean_absolute_error: 0.3474 - val_loss: 0.2472 -  
val_mean_absolute_error: 0.3554  
Epoch 31/100  
369/369 ━━━━━━━━━━━ 0s 1ms/step - loss: 0.2325 -  
mean_absolute_error: 0.3433 - val_loss: 0.2360 -  
val_mean_absolute_error: 0.3307  
Epoch 32/100  
369/369 ━━━━━━━━━━━ 0s 1ms/step - loss: 0.2322 -  
mean_absolute_error: 0.3397 - val_loss: 0.2346 -  
val_mean_absolute_error: 0.3375  
Epoch 33/100  
369/369 ━━━━━━━━━━━ 1s 2ms/step - loss: 0.2249 -  
mean_absolute_error: 0.3359 - val_loss: 0.2359 -  
val_mean_absolute_error: 0.3366  
Epoch 34/100  
369/369 ━━━━━━━━━━━ 1s 2ms/step - loss: 0.2291 -  
mean_absolute_error: 0.3394 - val_loss: 0.2313 -  
val_mean_absolute_error: 0.3330  
Epoch 35/100  
369/369 ━━━━━━━━━━━ 1s 1ms/step - loss: 0.2298 -  
mean_absolute_error: 0.3395 - val_loss: 0.2382 -
```

```
val_mean_absolute_error: 0.3366
Epoch 36/100
369/369 _____ 0s 1ms/step - loss: 0.2247 -
mean_absolute_error: 0.3327 - val_loss: 0.2355 -
val_mean_absolute_error: 0.3371
Epoch 37/100
369/369 _____ 0s 1ms/step - loss: 0.2211 -
mean_absolute_error: 0.3336 - val_loss: 0.2294 -
val_mean_absolute_error: 0.3316
Epoch 38/100
369/369 _____ 1s 1ms/step - loss: 0.2271 -
mean_absolute_error: 0.3370 - val_loss: 0.2277 -
val_mean_absolute_error: 0.3223
Epoch 39/100
369/369 _____ 1s 1ms/step - loss: 0.2266 -
mean_absolute_error: 0.3344 - val_loss: 0.2339 -
val_mean_absolute_error: 0.3335
Epoch 40/100
369/369 _____ 0s 1ms/step - loss: 0.2213 -
mean_absolute_error: 0.3302 - val_loss: 0.2313 -
val_mean_absolute_error: 0.3366
Epoch 41/100
369/369 _____ 0s 993us/step - loss: 0.2148 -
mean_absolute_error: 0.3311 - val_loss: 0.2292 -
val_mean_absolute_error: 0.3325
Epoch 42/100
369/369 _____ 0s 1ms/step - loss: 0.2246 -
mean_absolute_error: 0.3336 - val_loss: 0.2300 -
val_mean_absolute_error: 0.3234
Epoch 43/100
369/369 _____ 0s 1ms/step - loss: 0.2318 -
mean_absolute_error: 0.3370 - val_loss: 0.2310 -
val_mean_absolute_error: 0.3266
Epoch 44/100
369/369 _____ 0s 1ms/step - loss: 0.2152 -
mean_absolute_error: 0.3285 - val_loss: 0.2289 -
val_mean_absolute_error: 0.3303
Epoch 45/100
369/369 _____ 0s 1ms/step - loss: 0.2279 -
mean_absolute_error: 0.3371 - val_loss: 0.2346 -
val_mean_absolute_error: 0.3381
Epoch 46/100
369/369 _____ 0s 1ms/step - loss: 0.2278 -
mean_absolute_error: 0.3350 - val_loss: 0.2318 -
val_mean_absolute_error: 0.3344
Epoch 47/100
369/369 _____ 0s 1ms/step - loss: 0.2213 -
mean_absolute_error: 0.3323 - val_loss: 0.2266 -
val_mean_absolute_error: 0.3241
```

Epoch 48/100
369/369 _____ 0s 1ms/step - loss: 0.2158 -
mean_absolute_error: 0.3277 - val_loss: 0.2283 -
val_mean_absolute_error: 0.3308
Epoch 49/100
369/369 _____ 0s 1ms/step - loss: 0.2141 -
mean_absolute_error: 0.3262 - val_loss: 0.2329 -
val_mean_absolute_error: 0.3278
Epoch 50/100
369/369 _____ 0s 1ms/step - loss: 0.2111 -
mean_absolute_error: 0.3232 - val_loss: 0.2318 -
val_mean_absolute_error: 0.3270
Epoch 51/100
369/369 _____ 0s 1ms/step - loss: 0.2248 -
mean_absolute_error: 0.3338 - val_loss: 0.2241 -
val_mean_absolute_error: 0.3219
Epoch 52/100
369/369 _____ 0s 1ms/step - loss: 0.2216 -
mean_absolute_error: 0.3308 - val_loss: 0.2328 -
val_mean_absolute_error: 0.3350
Epoch 53/100
369/369 _____ 0s 1ms/step - loss: 0.2095 -
mean_absolute_error: 0.3242 - val_loss: 0.2338 -
val_mean_absolute_error: 0.3376
Epoch 54/100
369/369 _____ 0s 1ms/step - loss: 0.2174 -
mean_absolute_error: 0.3312 - val_loss: 0.2268 -
val_mean_absolute_error: 0.3248
Epoch 55/100
369/369 _____ 0s 1ms/step - loss: 0.2177 -
mean_absolute_error: 0.3280 - val_loss: 0.2277 -
val_mean_absolute_error: 0.3268
Epoch 56/100
369/369 _____ 0s 1ms/step - loss: 0.2083 -
mean_absolute_error: 0.3239 - val_loss: 0.2325 -
val_mean_absolute_error: 0.3301
Epoch 57/100
369/369 _____ 0s 1ms/step - loss: 0.2176 -
mean_absolute_error: 0.3278 - val_loss: 0.2272 -
val_mean_absolute_error: 0.3246
Epoch 58/100
369/369 _____ 0s 1ms/step - loss: 0.2226 -
mean_absolute_error: 0.3303 - val_loss: 0.2299 -
val_mean_absolute_error: 0.3236
Epoch 59/100
369/369 _____ 0s 1ms/step - loss: 0.2166 -
mean_absolute_error: 0.3276 - val_loss: 0.2266 -
val_mean_absolute_error: 0.3327
Epoch 60/100

```
369/369 _____ 0s 1ms/step - loss: 0.2162 -  
mean_absolute_error: 0.3271 - val_loss: 0.2271 -  
val_mean_absolute_error: 0.3289  
Epoch 61/100  
369/369 _____ 0s 1ms/step - loss: 0.2053 -  
mean_absolute_error: 0.3200 - val_loss: 0.2301 -  
val_mean_absolute_error: 0.3231  
Epoch 62/100  
369/369 _____ 0s 1ms/step - loss: 0.2098 -  
mean_absolute_error: 0.3229 - val_loss: 0.2377 -  
val_mean_absolute_error: 0.3465  
Epoch 63/100  
369/369 _____ 0s 1ms/step - loss: 0.2282 -  
mean_absolute_error: 0.3365 - val_loss: 0.2267 -  
val_mean_absolute_error: 0.3237  
Epoch 64/100  
369/369 _____ 0s 1ms/step - loss: 0.2114 -  
mean_absolute_error: 0.3216 - val_loss: 0.2256 -  
val_mean_absolute_error: 0.3315  
Epoch 65/100  
369/369 _____ 0s 1ms/step - loss: 0.2183 -  
mean_absolute_error: 0.3269 - val_loss: 0.2254 -  
val_mean_absolute_error: 0.3205  
Epoch 66/100  
369/369 _____ 0s 1ms/step - loss: 0.2142 -  
mean_absolute_error: 0.3258 - val_loss: 0.2283 -  
val_mean_absolute_error: 0.3257  
Epoch 67/100  
369/369 _____ 0s 1ms/step - loss: 0.2084 -  
mean_absolute_error: 0.3217 - val_loss: 0.2257 -  
val_mean_absolute_error: 0.3187  
Epoch 68/100  
369/369 _____ 0s 1ms/step - loss: 0.2073 -  
mean_absolute_error: 0.3199 - val_loss: 0.2240 -  
val_mean_absolute_error: 0.3180  
Epoch 69/100  
369/369 _____ 1s 1ms/step - loss: 0.2128 -  
mean_absolute_error: 0.3250 - val_loss: 0.2285 -  
val_mean_absolute_error: 0.3256  
Epoch 70/100  
369/369 _____ 0s 1ms/step - loss: 0.2130 -  
mean_absolute_error: 0.3249 - val_loss: 0.2270 -  
val_mean_absolute_error: 0.3213  
Epoch 71/100  
369/369 _____ 0s 1ms/step - loss: 0.2127 -  
mean_absolute_error: 0.3240 - val_loss: 0.2255 -  
val_mean_absolute_error: 0.3188  
Epoch 72/100  
369/369 _____ 0s 1ms/step - loss: 0.2023 -
```



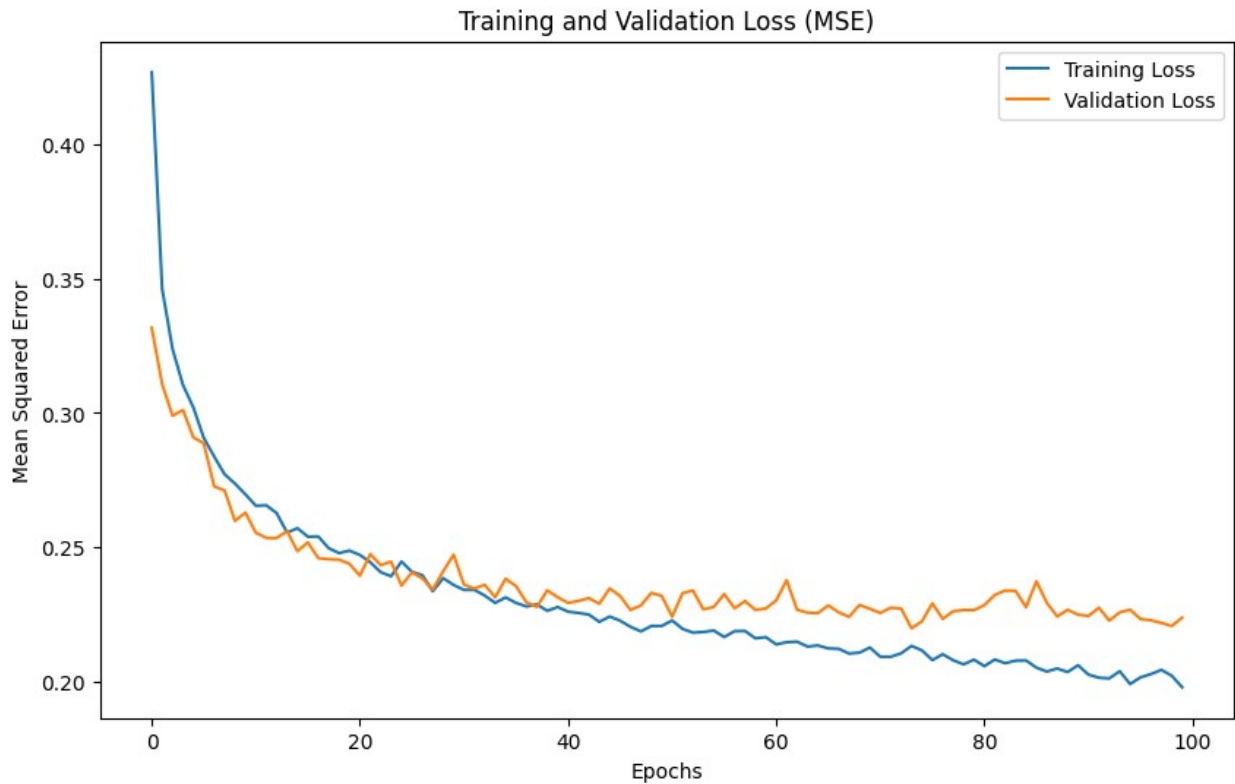
```
mean_absolute_error: 0.3163 - val_loss: 0.2274 -  
val_mean_absolute_error: 0.3192  
Epoch 73/100  
369/369 _____ 1s 1ms/step - loss: 0.2071 -  
mean_absolute_error: 0.3193 - val_loss: 0.2271 -  
val_mean_absolute_error: 0.3275  
Epoch 74/100  
369/369 _____ 0s 1ms/step - loss: 0.2139 -  
mean_absolute_error: 0.3253 - val_loss: 0.2198 -  
val_mean_absolute_error: 0.3169  
Epoch 75/100  
369/369 _____ 0s 1ms/step - loss: 0.2183 -  
mean_absolute_error: 0.3281 - val_loss: 0.2223 -  
val_mean_absolute_error: 0.3219  
Epoch 76/100  
369/369 _____ 0s 1ms/step - loss: 0.2072 -  
mean_absolute_error: 0.3192 - val_loss: 0.2291 -  
val_mean_absolute_error: 0.3283  
Epoch 77/100  
369/369 _____ 0s 1ms/step - loss: 0.2083 -  
mean_absolute_error: 0.3177 - val_loss: 0.2233 -  
val_mean_absolute_error: 0.3174  
Epoch 78/100  
369/369 _____ 0s 1ms/step - loss: 0.2125 -  
mean_absolute_error: 0.3211 - val_loss: 0.2261 -  
val_mean_absolute_error: 0.3224  
Epoch 79/100  
369/369 _____ 0s 1ms/step - loss: 0.1989 -  
mean_absolute_error: 0.3143 - val_loss: 0.2266 -  
val_mean_absolute_error: 0.3201  
Epoch 80/100  
369/369 _____ 0s 1ms/step - loss: 0.2142 -  
mean_absolute_error: 0.3219 - val_loss: 0.2265 -  
val_mean_absolute_error: 0.3310  
Epoch 81/100  
369/369 _____ 0s 1ms/step - loss: 0.2110 -  
mean_absolute_error: 0.3240 - val_loss: 0.2284 -  
val_mean_absolute_error: 0.3315  
Epoch 82/100  
369/369 _____ 1s 1ms/step - loss: 0.2103 -  
mean_absolute_error: 0.3221 - val_loss: 0.2322 -  
val_mean_absolute_error: 0.3293  
Epoch 83/100  
369/369 _____ 1s 1ms/step - loss: 0.2000 -  
mean_absolute_error: 0.3143 - val_loss: 0.2338 -  
val_mean_absolute_error: 0.3385  
Epoch 84/100  
369/369 _____ 0s 1ms/step - loss: 0.2034 -  
mean_absolute_error: 0.3196 - val_loss: 0.2337 -
```

```
val_mean_absolute_error: 0.3224
Epoch 85/100
369/369 _____ 0s 1ms/step - loss: 0.2028 -
mean_absolute_error: 0.3204 - val_loss: 0.2277 -
val_mean_absolute_error: 0.3258
Epoch 86/100
369/369 _____ 0s 1ms/step - loss: 0.1974 -
mean_absolute_error: 0.3151 - val_loss: 0.2373 -
val_mean_absolute_error: 0.3410
Epoch 87/100
369/369 _____ 0s 1ms/step - loss: 0.1960 -
mean_absolute_error: 0.3138 - val_loss: 0.2292 -
val_mean_absolute_error: 0.3204
Epoch 88/100
369/369 _____ 0s 1ms/step - loss: 0.2029 -
mean_absolute_error: 0.3148 - val_loss: 0.2242 -
val_mean_absolute_error: 0.3181
Epoch 89/100
369/369 _____ 0s 1ms/step - loss: 0.2018 -
mean_absolute_error: 0.3156 - val_loss: 0.2266 -
val_mean_absolute_error: 0.3280
Epoch 90/100
369/369 _____ 0s 1ms/step - loss: 0.2000 -
mean_absolute_error: 0.3156 - val_loss: 0.2249 -
val_mean_absolute_error: 0.3210
Epoch 91/100
369/369 _____ 0s 1ms/step - loss: 0.1884 -
mean_absolute_error: 0.3071 - val_loss: 0.2243 -
val_mean_absolute_error: 0.3226
Epoch 92/100
369/369 _____ 0s 1ms/step - loss: 0.2035 -
mean_absolute_error: 0.3212 - val_loss: 0.2274 -
val_mean_absolute_error: 0.3282
Epoch 93/100
369/369 _____ 0s 1ms/step - loss: 0.1982 -
mean_absolute_error: 0.3155 - val_loss: 0.2226 -
val_mean_absolute_error: 0.3186
Epoch 94/100
369/369 _____ 0s 1ms/step - loss: 0.2045 -
mean_absolute_error: 0.3185 - val_loss: 0.2257 -
val_mean_absolute_error: 0.3245
Epoch 95/100
369/369 _____ 0s 1ms/step - loss: 0.1955 -
mean_absolute_error: 0.3154 - val_loss: 0.2267 -
val_mean_absolute_error: 0.3251
Epoch 96/100
369/369 _____ 0s 1ms/step - loss: 0.2055 -
mean_absolute_error: 0.3183 - val_loss: 0.2232 -
val_mean_absolute_error: 0.3238
```

```
Epoch 97/100
369/369 _____ 0s 1ms/step - loss: 0.2033 -
mean_absolute_error: 0.3175 - val_loss: 0.2227 -
val_mean_absolute_error: 0.3266
Epoch 98/100
369/369 _____ 0s 1ms/step - loss: 0.2014 -
mean_absolute_error: 0.3172 - val_loss: 0.2218 -
val_mean_absolute_error: 0.3221
Epoch 99/100
369/369 _____ 0s 1ms/step - loss: 0.1963 -
mean_absolute_error: 0.3102 - val_loss: 0.2206 -
val_mean_absolute_error: 0.3162
Epoch 100/100
369/369 _____ 0s 1ms/step - loss: 0.1940 -
mean_absolute_error: 0.3110 - val_loss: 0.2237 -
val_mean_absolute_error: 0.3183
```

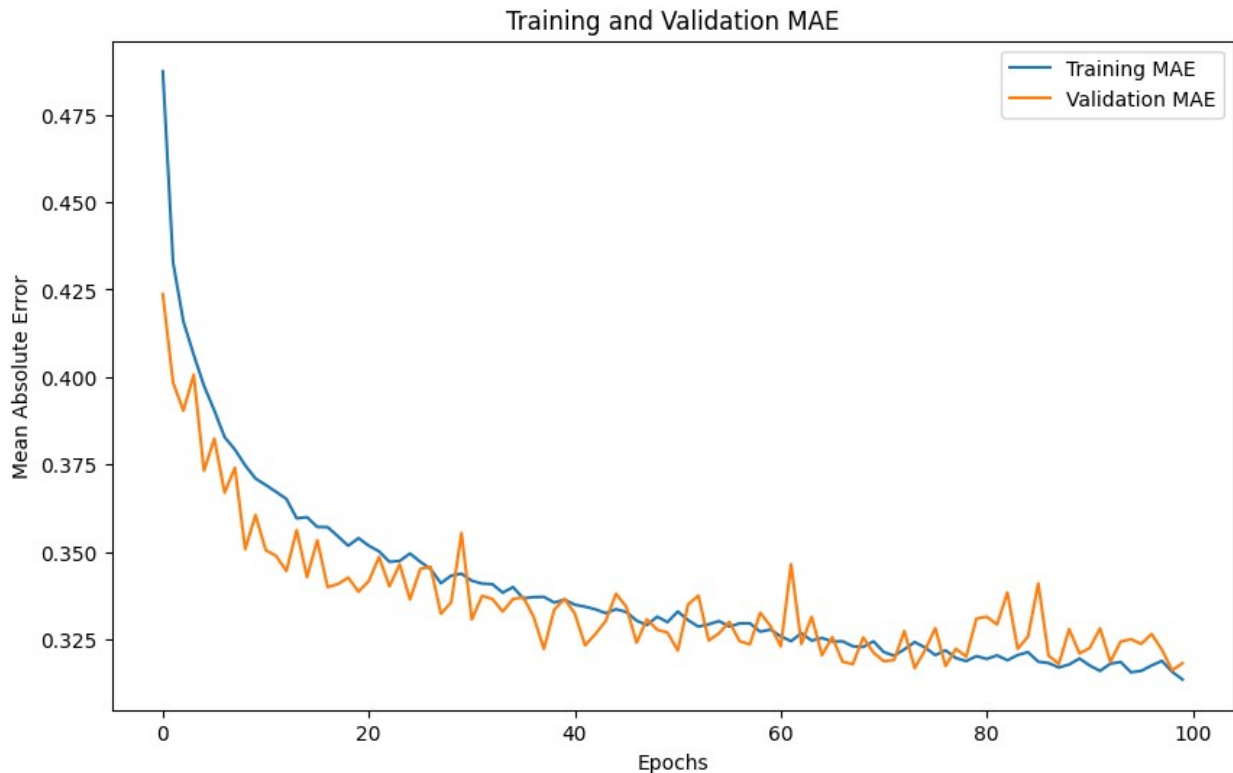
Let's plot the training history.

```
plt.figure(figsize=(10, 6))
plt.plot(history.history['loss'], label='Training Loss')
plt.plot(history.history['val_loss'], label='Validation Loss')
plt.xlabel('Epochs')
plt.ylabel('Mean Squared Error')
plt.title('Training and Validation Loss (MSE)')
plt.legend()
plt.show()
```



Let's also plot how the MAE change during training.

```
plt.figure(figsize=(10, 6))
plt.plot(history.history['mean_absolute_error'], label='Training MAE')
plt.plot(history.history['val_mean_absolute_error'], label='Validation MAE')
plt.xlabel('Epochs')
plt.ylabel('Mean Absolute Error')
plt.title('Training and Validation MAE')
plt.legend()
plt.show()
```



5. Model Evaluation

```
mse, mae = model.evaluate(X_test, y_test)
print(f"Mean Squared Error on test set: {mse}")
print(f"Mean Absolute Error on test set: {mae}")
```

```
123/123 ————— 0s 772us/step - loss: 0.2306 -
mean_absolute_error: 0.3253
Mean Squared Error on test set: 0.2286626100540161
Mean Absolute Error on test set: 0.325112909078598
```

Oops, we need to revert back the scaling to get the true value of the predicted median_housing_value.

```
# Inverse transform the predictions and the true values
y_pred_scaled = model.predict(X_test)
y_pred = scaler.inverse_transform(y_pred_scaled).flatten()
y_test_original = scaler.inverse_transform(y_test.reshape(-1,
1)).flatten()

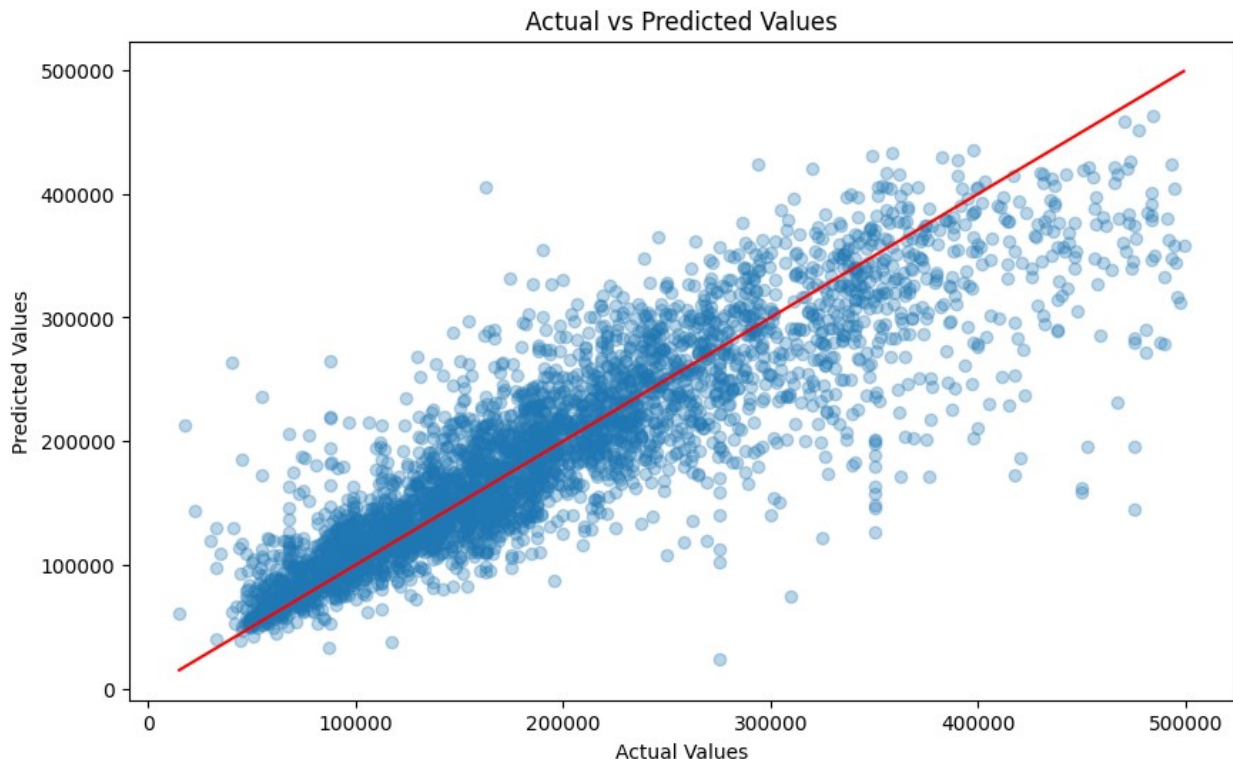
# Calculate the Mean Squared Error and Mean Absolute Error in the
original scale
mse_original = np.mean((y_pred - y_test_original) ** 2)
mae_original = np.mean(np.abs(y_pred - y_test_original))

print(f"Mean Squared Error on test set (original scale):
```

```
{mse_original}")
print(f"Mean Absolute Error on test set (original scale):
{mae_original}")

# Plot predictions vs actual values
plt.figure(figsize=(10, 6))
plt.scatter(y_test_original, y_pred, alpha=0.3)
plt.plot([min(y_test_original), max(y_test_original)],
[min(y_test_original), max(y_test_original)], color='red')
plt.xlabel('Actual Values')
plt.ylabel('Predicted Values')
plt.title('Actual vs Predicted Values')
plt.show()

123/123 ————— 0s 782us/step
Mean Squared Error on test set (original scale): 2184271616.0
Mean Absolute Error on test set (original scale): 31775.314453125
```



Observations from MAE of testing prediction:

- MAE = 31,775.
- It is quite accurate to be honest.

Observations from Actual vs Predicted plot:

- The points are more concentrated along the red line in the lower value range (below 200,000), which indicates better performance for houses in this price range.

- The spread of points increases as the actual value increases, which shows that the model is less accurate for higher-priced houses.
- The points appear to be more dispersed for actual values above 300,000, indicating higher prediction errors in this range.