

fd001-eda

June 14, 2024

1 Predictive Maintenance of Turbofan Jet Engine: Exploratory Data Analysis

1.1 1. Data Inspection

```
[1]: import warnings
warnings.filterwarnings("ignore")

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
[2]: # Open dataset description README.md
with open("data/README.md", "r") as f:
    print(f.read())
```

Data Set: FD001
Train trajectories: 100
Test trajectories: 100
Conditions: ONE (Sea Level)
Fault Modes: ONE (HPC Degradation)

Data Set: FD002
Train trajectories: 260
Test trajectories: 259
Conditions: SIX
Fault Modes: ONE (HPC Degradation)

Data Set: FD003
Train trajectories: 100
Test trajectories: 100
Conditions: ONE (Sea Level)
Fault Modes: TWO (HPC Degradation, Fan Degradation)

Data Set: FD004
Train trajectories: 248
Test trajectories: 249

Conditions: SIX

Fault Modes: TWO (HPC Degradation, Fan Degradation)

Experimental Scenario

Data sets consists of multiple multivariate time series. Each data set is further divided into training and test subsets. Each time series is from a different engine i.e., the data can be considered to be from a fleet of engines of the same type. Each engine starts with different degrees of initial wear and manufacturing variation which is unknown to the user. This wear and variation is considered normal, i.e., it is not considered a fault condition. There are three operational settings that have a substantial effect on engine performance. These settings are also included in the data. The data is contaminated with sensor noise.

The engine is operating normally at the start of each time series, and develops a fault at some point during the series. In the training set, the fault grows in magnitude until system failure. In the test set, the time series ends some time prior to system failure. The objective of the competition is to predict the number of remaining operational cycles before failure in the test set, i.e., the number of operational cycles after the last cycle that the engine will continue to operate. Also provided a vector of true Remaining Useful Life (RUL) values for the test data.

The data are provided as a zip-compressed text file with 26 columns of numbers, separated by spaces. Each row is a snapshot of data taken during a single operational cycle, each column is a different variable. The columns correspond to:

- 1) unit number
- 2) time, in cycles
- 3) operational setting 1
- 4) operational setting 2
- 5) operational setting 3
- 6) sensor measurement 1
- 7) sensor measurement 2
- ...
- 26) sensor measurement 21

Reference: A. Saxena, K. Goebel, D. Simon, and N. Eklund, iDamage Propagation Modeling for Aircraft Engine Run-to-Failure Simulation, in the Proceedings of the 1st International Conference on Prognostics and Health Management (PHM08), Denver CO, Oct 2008.

For this assignment, we will be using **FD001** dataset since it is widely researched. Thus, it is

suitable for benchmarking purposes.

```
[3]: # Define the columns of our dataset.
INDEX_COLUMNS = ["unit", "time_cycles"]
OP_SETTING_COLUMNS = ["op_setting_{}".format(i) for i in range(1, 4)]
SENSOR_COLUMNS = ["sensor_{}".format(i) for i in range(1, 22)]
col_names = INDEX_COLUMNS + OP_SETTING_COLUMNS + SENSOR_COLUMNS
```

```
[4]: # Load training data
train_data = pd.read_csv(
    "data/train_FD001.txt", sep="\s+", header=None, names=col_names
)
train_data.head()
```

```
[4]:   unit  time_cycles  op_setting_1  op_setting_2  op_setting_3  sensor_1  \
0     1             1      -0.0007      -0.0004           100.0    518.67
1     1             2       0.0019      -0.0003           100.0    518.67
2     1             3      -0.0043       0.0003           100.0    518.67
3     1             4       0.0007       0.0000           100.0    518.67
4     1             5      -0.0019      -0.0002           100.0    518.67
```

```
      sensor_2  sensor_3  sensor_4  sensor_5  ...  sensor_12  sensor_13  \
0      641.82   1589.70   1400.60    14.62  ...     521.66    2388.02
1      642.15   1591.82   1403.14    14.62  ...     522.28    2388.07
2      642.35   1587.99   1404.20    14.62  ...     522.42    2388.03
3      642.35   1582.79   1401.87    14.62  ...     522.86    2388.08
4      642.37   1582.85   1406.22    14.62  ...     522.19    2388.04
```

```
      sensor_14  sensor_15  sensor_16  sensor_17  sensor_18  sensor_19  \
0      8138.62     8.4195     0.03         392         2388         100.0
1      8131.49     8.4318     0.03         392         2388         100.0
2      8133.23     8.4178     0.03         390         2388         100.0
3      8133.83     8.3682     0.03         392         2388         100.0
4      8133.80     8.4294     0.03         393         2388         100.0
```

```
      sensor_20  sensor_21
0         39.06    23.4190
1         39.00    23.4236
2         38.95    23.3442
3         38.88    23.3739
4         38.90    23.4044
```

[5 rows x 26 columns]

```
[5]: # Load test data
test_data = pd.read_csv("data/test_FD001.txt", sep="\s+", header=None,
    ↪ names=col_names)
```

```
test_data.head()
```

```
[5]:   unit  time_cycles  op_setting_1  op_setting_2  op_setting_3  sensor_1  \
0     1             1         0.0023         0.0003          100.0    518.67
1     1             2        -0.0027        -0.0003          100.0    518.67
2     1             3         0.0003         0.0001          100.0    518.67
3     1             4         0.0042         0.0000          100.0    518.67
4     1             5         0.0014         0.0000          100.0    518.67

      sensor_2  sensor_3  sensor_4  sensor_5  ...  sensor_12  sensor_13  \
0      643.02   1585.29   1398.21    14.62  ...     521.72    2388.03
1      641.71   1588.45   1395.42    14.62  ...     522.16    2388.06
2      642.46   1586.94   1401.34    14.62  ...     521.97    2388.03
3      642.44   1584.12   1406.42    14.62  ...     521.38    2388.05
4      642.51   1587.19   1401.92    14.62  ...     522.15    2388.03

      sensor_14  sensor_15  sensor_16  sensor_17  sensor_18  sensor_19  \
0      8125.55     8.4052     0.03         392         2388         100.0
1      8139.62     8.3803     0.03         393         2388         100.0
2      8130.10     8.4441     0.03         393         2388         100.0
3      8132.90     8.3917     0.03         391         2388         100.0
4      8129.54     8.4031     0.03         390         2388         100.0

      sensor_20  sensor_21
0         38.86    23.3735
1         39.02    23.3916
2         39.08    23.4166
3         39.00    23.3737
4         38.99    23.4130
```

[5 rows x 26 columns]

```
[6]: # Load the true RUL values for the test data
test_rul = pd.read_csv("data/RUL_FD001.txt", sep="\s+", header=None,
    names=["rul"])

# Add a unit column to the test RUL
test_rul["unit"] = test_rul.index + 1
test_rul.head()
```

```
[6]:   rul  unit
0  112     1
1   98     2
2   69     3
3   82     4
4   91     5
```

```
[7]: train_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20631 entries, 0 to 20630
Data columns (total 26 columns):
 #   Column          Non-Null Count  Dtype
---  -
 0   unit            20631 non-null  int64
 1   time_cycles     20631 non-null  int64
 2   op_setting_1    20631 non-null  float64
 3   op_setting_2    20631 non-null  float64
 4   op_setting_3    20631 non-null  float64
 5   sensor_1        20631 non-null  float64
 6   sensor_2        20631 non-null  float64
 7   sensor_3        20631 non-null  float64
 8   sensor_4        20631 non-null  float64
 9   sensor_5        20631 non-null  float64
10   sensor_6        20631 non-null  float64
11   sensor_7        20631 non-null  float64
12   sensor_8        20631 non-null  float64
13   sensor_9        20631 non-null  float64
14   sensor_10       20631 non-null  float64
15   sensor_11       20631 non-null  float64
16   sensor_12       20631 non-null  float64
17   sensor_13       20631 non-null  float64
18   sensor_14       20631 non-null  float64
19   sensor_15       20631 non-null  float64
20   sensor_16       20631 non-null  float64
21   sensor_17       20631 non-null  int64
22   sensor_18       20631 non-null  int64
23   sensor_19       20631 non-null  float64
24   sensor_20       20631 non-null  float64
25   sensor_21       20631 non-null  float64
dtypes: float64(22), int64(4)
memory usage: 4.1 MB
```

```
[8]: train_data.describe()
```

```
[8]:
```

	unit	time_cycles	op_setting_1	op_setting_2	op_setting_3	\
count	20631.000000	20631.000000	20631.000000	20631.000000	20631.0	
mean	51.506568	108.807862	-0.000009	0.000002	100.0	
std	29.227633	68.880990	0.002187	0.000293	0.0	
min	1.000000	1.000000	-0.008700	-0.000600	100.0	
25%	26.000000	52.000000	-0.001500	-0.000200	100.0	
50%	52.000000	104.000000	0.000000	0.000000	100.0	
75%	77.000000	156.000000	0.001500	0.000300	100.0	
max	100.000000	362.000000	0.008700	0.000600	100.0	

	sensor_1	sensor_2	sensor_3	sensor_4	sensor_5	...	\
count	20631.00	20631.000000	20631.000000	20631.000000	2.063100e+04	...	
mean	518.67	642.680934	1590.523119	1408.933782	1.462000e+01	...	
std	0.00	0.500053	6.131150	9.000605	1.776400e-15	...	
min	518.67	641.210000	1571.040000	1382.250000	1.462000e+01	...	
25%	518.67	642.325000	1586.260000	1402.360000	1.462000e+01	...	
50%	518.67	642.640000	1590.100000	1408.040000	1.462000e+01	...	
75%	518.67	643.000000	1594.380000	1414.555000	1.462000e+01	...	
max	518.67	644.530000	1616.910000	1441.490000	1.462000e+01	...	

	sensor_12	sensor_13	sensor_14	sensor_15	sensor_16	...	\
count	20631.000000	20631.000000	20631.000000	20631.000000	2.063100e+04	...	
mean	521.413470	2388.096152	8143.752722	8.442146	3.000000e-02	...	
std	0.737553	0.071919	19.076176	0.037505	1.387812e-17	...	
min	518.690000	2387.880000	8099.940000	8.324900	3.000000e-02	...	
25%	520.960000	2388.040000	8133.245000	8.414900	3.000000e-02	...	
50%	521.480000	2388.090000	8140.540000	8.438900	3.000000e-02	...	
75%	521.950000	2388.140000	8148.310000	8.465600	3.000000e-02	...	
max	523.380000	2388.560000	8293.720000	8.584800	3.000000e-02	...	

	sensor_17	sensor_18	sensor_19	sensor_20	sensor_21
count	20631.000000	20631.0	20631.0	20631.000000	20631.000000
mean	393.210654	2388.0	100.0	38.816271	23.289705
std	1.548763	0.0	0.0	0.180746	0.108251
min	388.000000	2388.0	100.0	38.140000	22.894200
25%	392.000000	2388.0	100.0	38.700000	23.221800
50%	393.000000	2388.0	100.0	38.830000	23.297900
75%	394.000000	2388.0	100.0	38.950000	23.366800
max	400.000000	2388.0	100.0	39.430000	23.618400

[8 rows x 26 columns]

1.2 2. Distribution of Engine Lifetime in Train Sets and Last Time Cycle in Test Sets

```
[9]: # Find the lifetime of each engine unit
train_units_lifetime = train_data.groupby("unit")["time_cycles"].max().
    ↪reset_index()
test_units_lifetime = test_data.groupby("unit")["time_cycles"].max().
    ↪reset_index()

# Add a column to indicate if the data is from the training or test set
train_units_lifetime["dataset"] = "train"
test_units_lifetime["dataset"] = "test"
```

```
units_lifetime = train_units_lifetime._append(test_units_lifetime, ignore_index=True)
```

```
[10]: units_lifetime.groupby("dataset")["time_cycles"].describe()
```

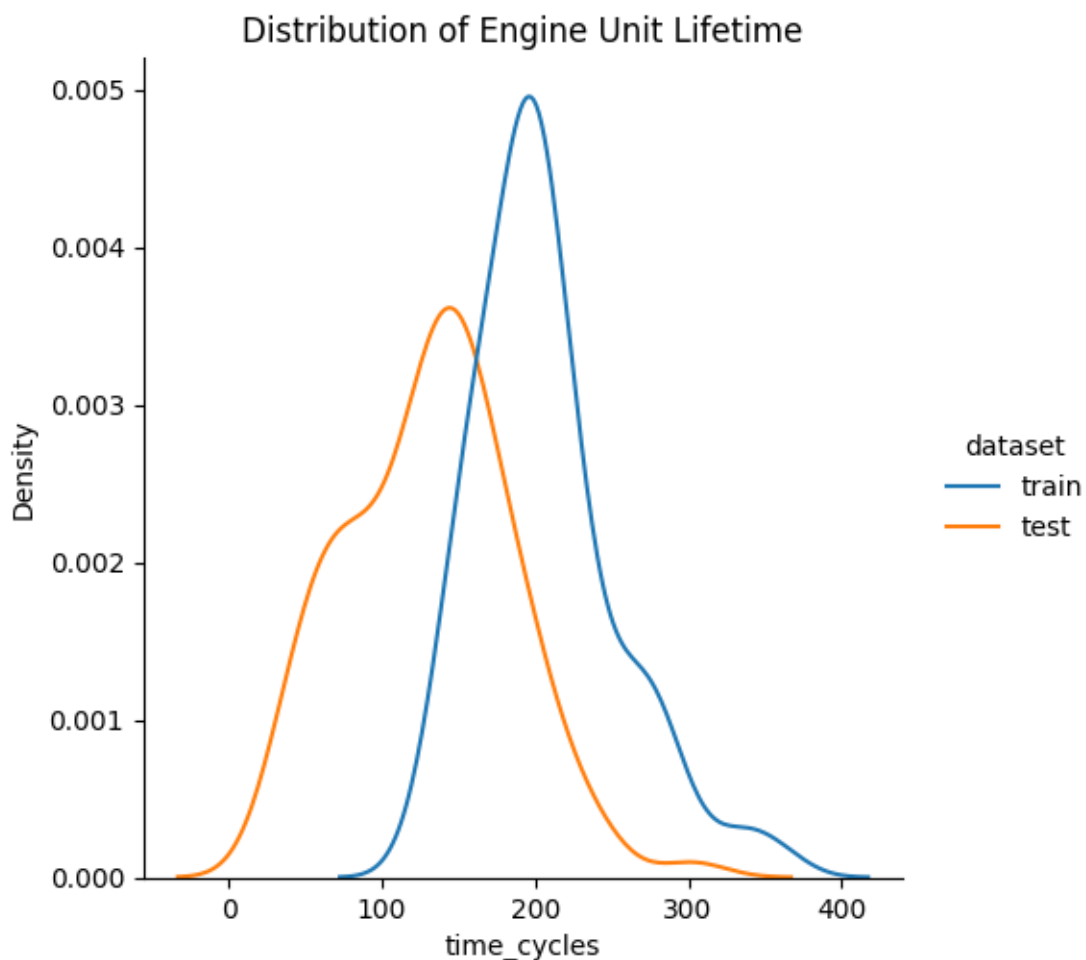
```
[10]:
```

	count	mean	std	min	25%	50%	75%	max
dataset								
test	100.0	130.96	53.593479	31.0	88.75	133.5	164.25	303.0
train	100.0	206.31	46.342749	128.0	177.00	199.0	229.25	362.0

```
[11]: # Plot the distribution of the lifetime of the engine units
plt.figure(figsize=(8, 6))
sns.displot(units_lifetime, x="time_cycles", hue="dataset", kind="kde")
plt.title("Distribution of Engine Unit Lifetime")
```

```
[11]: Text(0.5, 1.0, 'Distribution of Engine Unit Lifetime')
```

<Figure size 800x600 with 0 Axes>



Observations:

- Engine units in the test set have shorter lifetimes than train set.
- On average, engine lifetimes in the test set are 70 cycles shorter than train set.

1.3 3. Distribution of RUL in Train and Test Sets

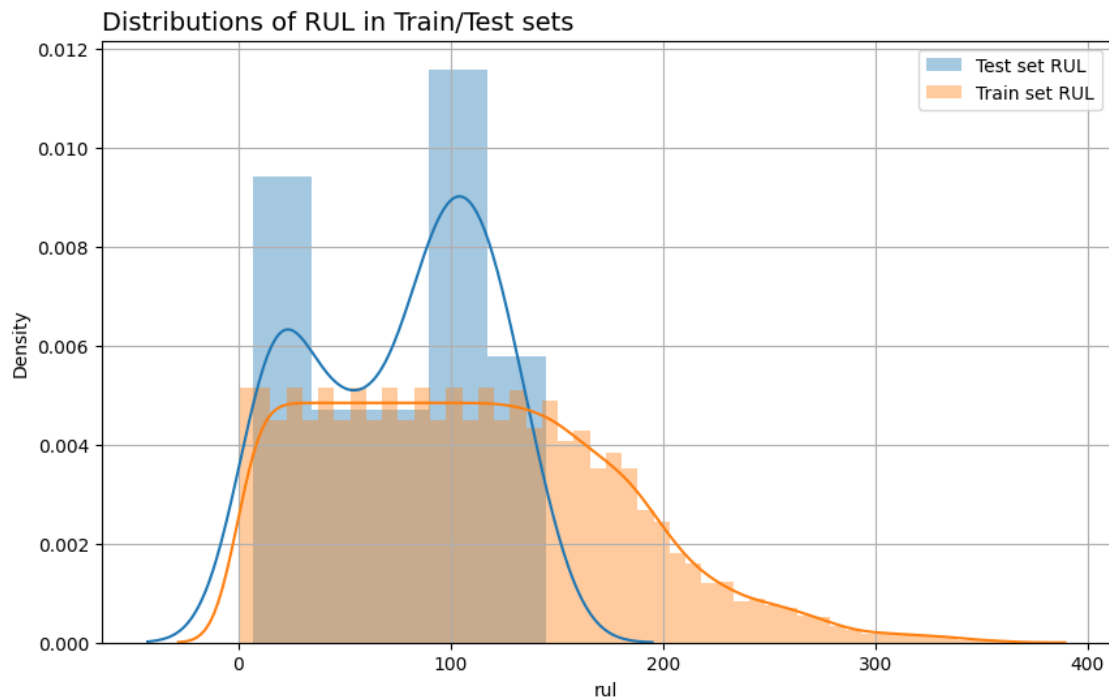
```
[12]: def calculate_RUL(X, upper_threshold=None):
        lifetime = X.groupby(["unit"])["time_cycles"].transform(max)
        rul = lifetime - X["time_cycles"]

        if upper_threshold:
            rul = np.where(rul > upper_threshold, upper_threshold, rul)

        return rul
```

```
[13]: train_data["rul"] = calculate_RUL(train_data)

fig, ax = plt.subplots(figsize=(10, 6))
ax.set_title("Distributions of RUL in Train/Test sets", loc="left", size=14)
sns.distplot(test_rul["rul"], label="Test set RUL")
sns.distplot(train_data["rul"], label="Train set RUL")
ax.legend()
ax.grid()
plt.show()
```



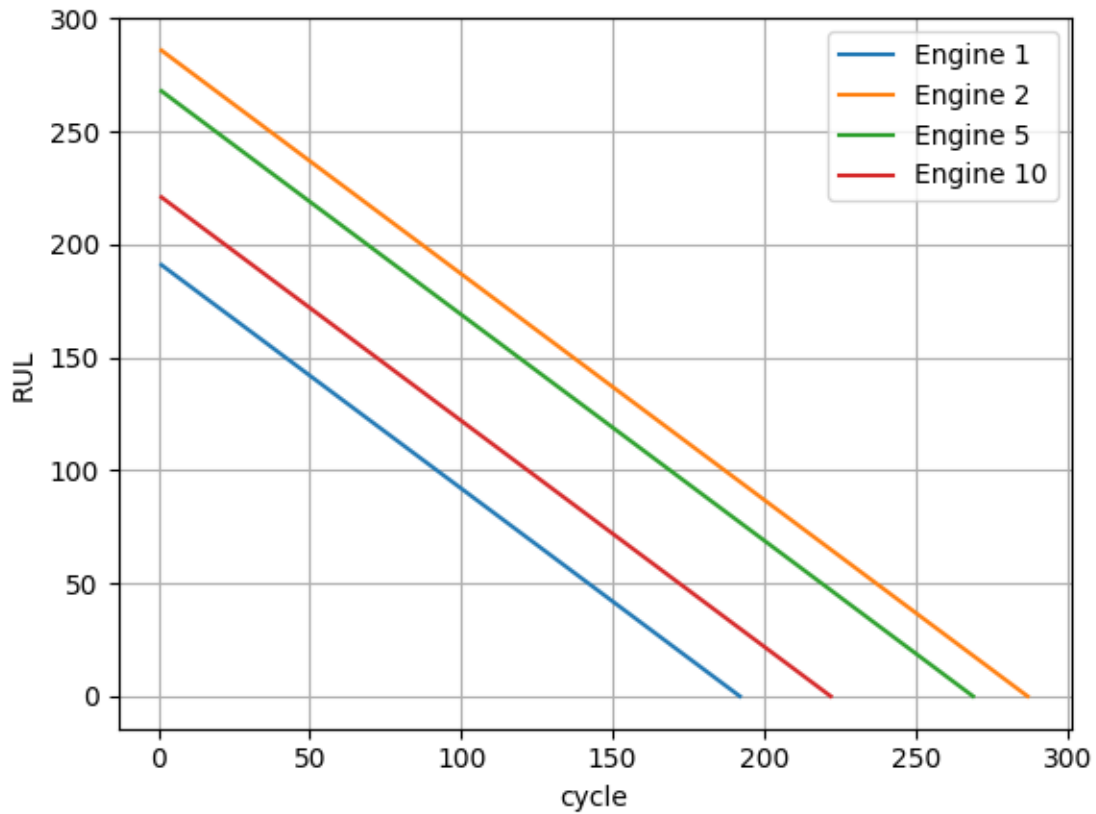
Observations:

- As expected, the Train set contains units with RUL way higher than in Test set.
- This is because the Train set have the full lifetime data of an engine while the Test set only contains data until a period of time before its end of life.
- Hence, our target is to predict the RUL at the last cycle of each engine in the Test set.

To address this issue, we need to limit the maximum value of RUL. The motivation is that a degradation process will only be noticeable in the data after a unit has operated for some time.

```
[14]: # Before applying limit
for _unit in [1, 2, 5, 10]:
    plt.plot(
        train_data[train_data["unit"] == _unit]["time_cycles"],
        train_data[train_data["unit"] == _unit]["rul"],
        label=f"Engine {_unit}",
    )

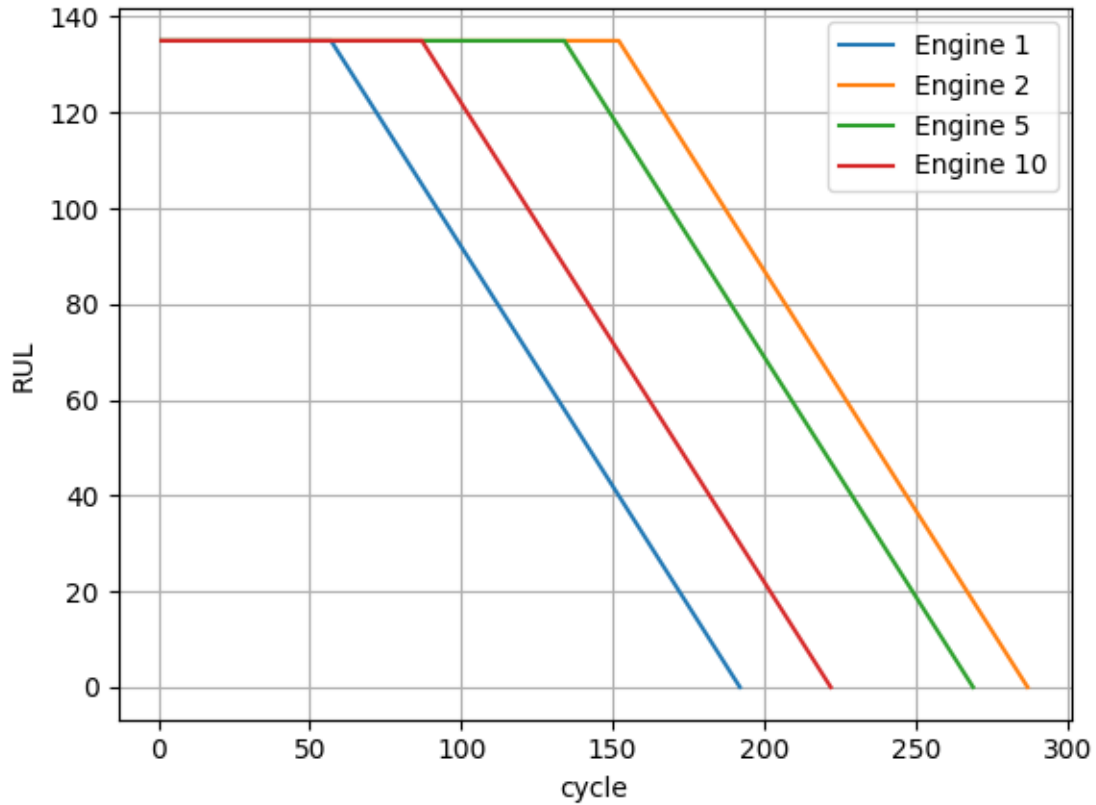
plt.legend()
plt.xlabel("cycle")
plt.ylabel("RUL")
plt.grid()
plt.show()
```



```
[15]: # After applying limit
train_data["rul"] = calculate_RUL(train_data, upper_threshold=135)

for _unit in [1, 2, 5, 10]:
    plt.plot(
        train_data[train_data["unit"] == _unit]["time_cycles"],
        train_data[train_data["unit"] == _unit]["rul"],
        label=f"Engine {_unit}",
    )

plt.legend()
plt.xlabel("cycle")
plt.ylabel("RUL")
plt.grid()
plt.show()
```

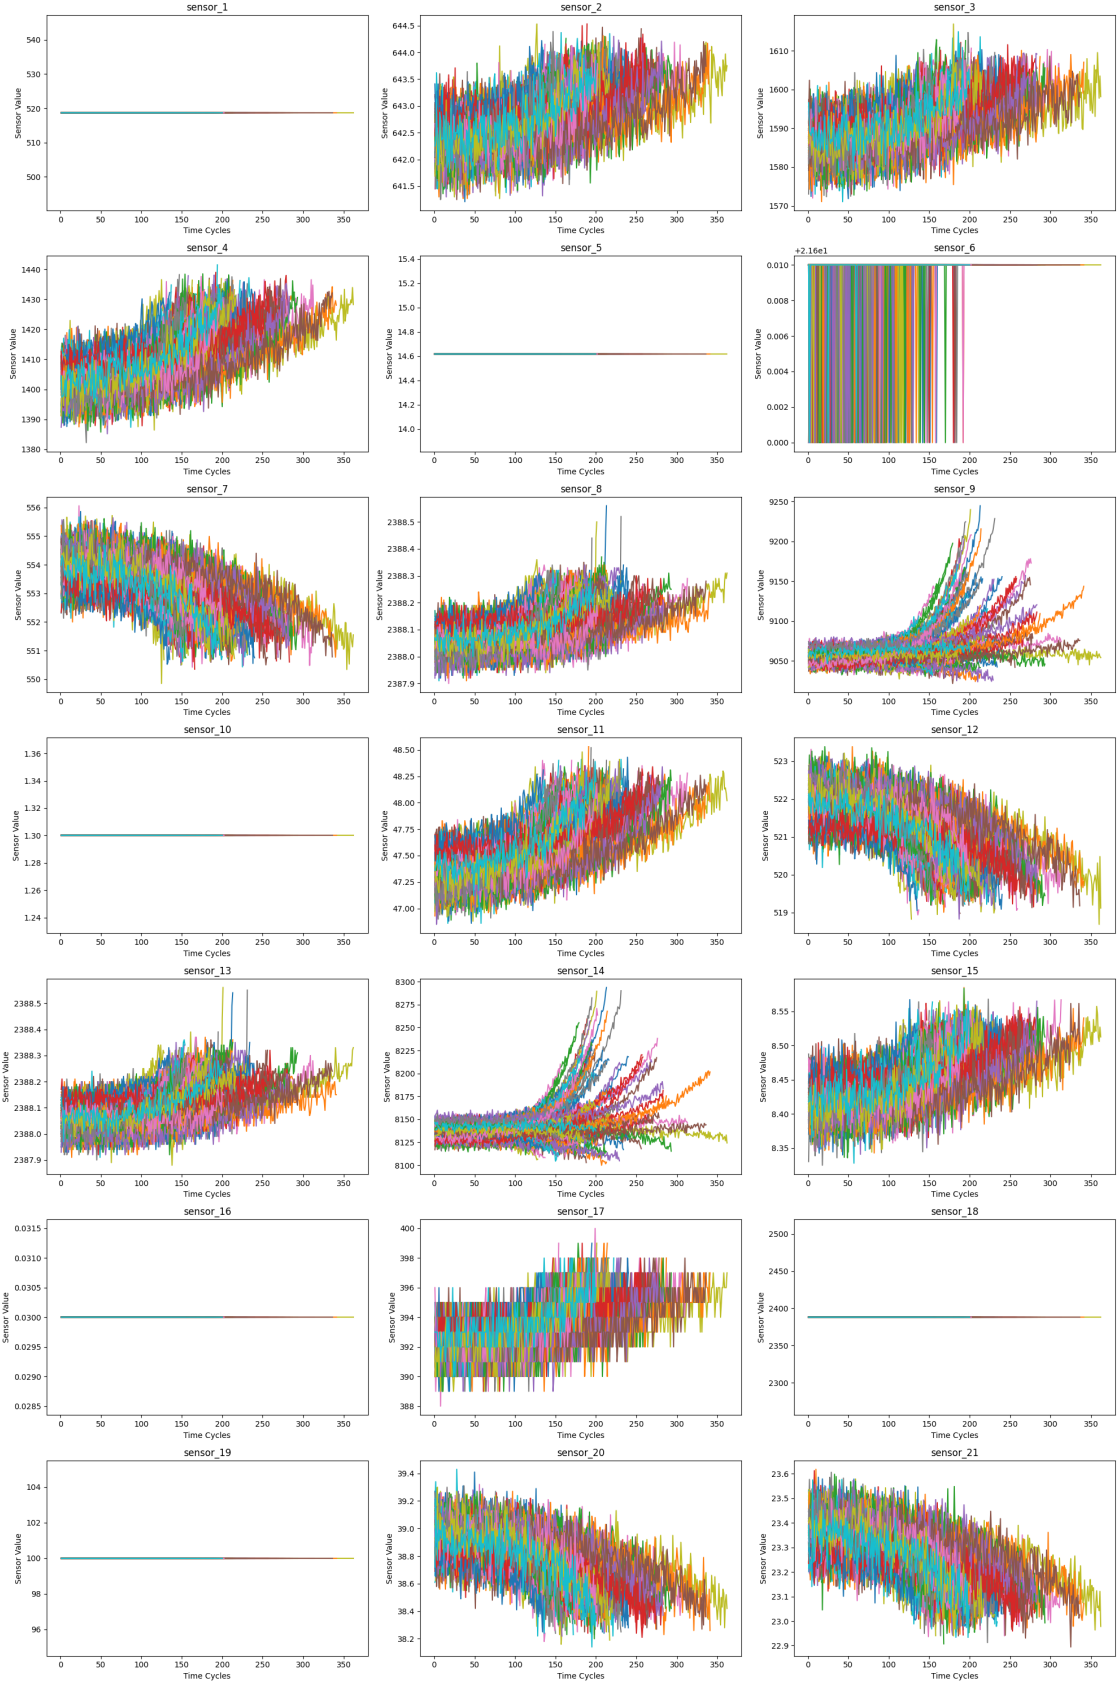


1.4 4. Wear / Degradation Patterns

```
[16]: fig, axes = plt.subplots(7, 3, figsize=(20, 30))
      axes = axes.flatten()

      for i, col in enumerate(SENSOR_COLUMNS):
          for unit in train_data["unit"].unique():
              unit_data = train_data[train_data["unit"] == unit]
              axes[i].plot(unit_data["time_cycles"], unit_data[col])
      axes[i].set_title(SENSOR_COLUMNS[i])
      axes[i].set_xlabel("Time Cycles")
      axes[i].set_ylabel("Sensor Value")

      plt.tight_layout()
      plt.show()
```



Observations:

- Sensor 1, 5, 10, 16, 18, 19 are constant throughout cycles.
- Due to low variance, they are considered useless features and can be dropped.
- For others, we can really see how the sensor values changed with respect to time cycles. Some increase and some decrease.

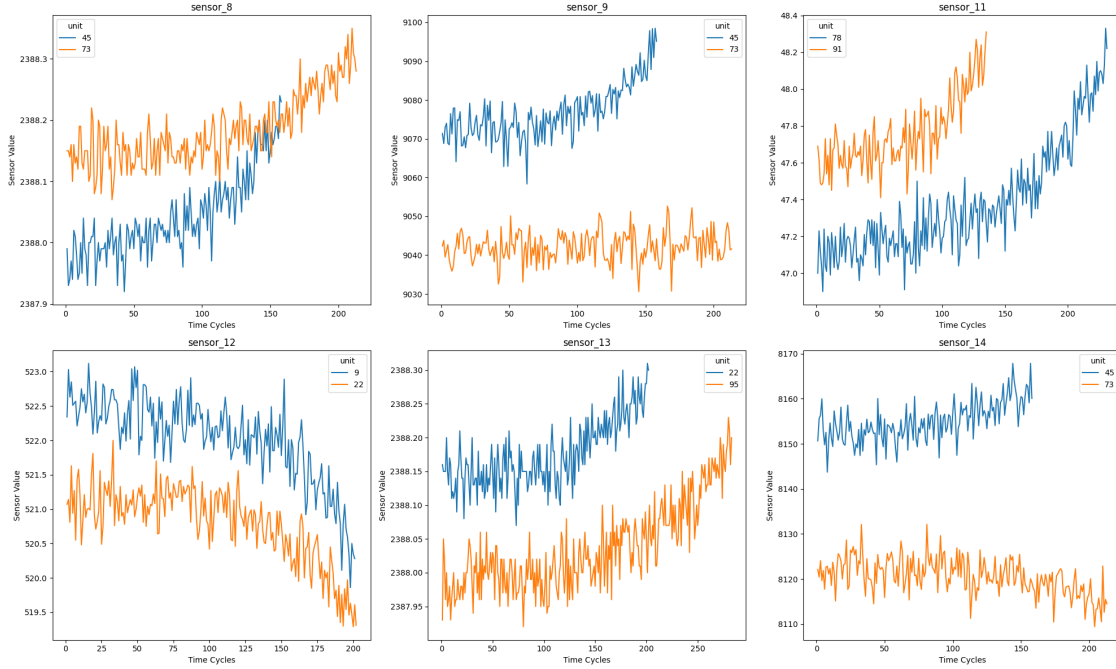
1.5 5. Difference between Each Unit of Engines

```
[17]: SELECTED_SENSORS = [
        "sensor_8",
        "sensor_9",
        "sensor_11",
        "sensor_12",
        "sensor_13",
        "sensor_14",
    ]

    first_avg_values = (
        train_data[train_data["time_cycles"] <= 20]
        .groupby("unit")[SELECTED_SENSORS]
        .mean()
        .reset_index()
    )

    fig, ax = plt.subplots(2, 3, figsize=(20, 12))
    for i, sensor in enumerate(SELECTED_SENSORS):
        avg_sorted = first_avg_values.sort_values(sensor)
        engine_a, engine_b = avg_sorted["unit"].iloc[0], avg_sorted["unit"].iloc[-1]
        sns.lineplot(
            data=train_data[train_data["unit"].isin([engine_a, engine_b])],
            x="time_cycles",
            y=sensor,
            hue="unit",
            ax=ax.flatten()[i],
            palette="tab10",
        )
        ax.flatten()[i].set_title(sensor)
        ax.flatten()[i].set_xlabel("Time Cycles")
        ax.flatten()[i].set_ylabel("Sensor Value")

    plt.tight_layout()
    plt.show()
```



Observations:

- We can see the difference between engines based on the sensor time series.
- This suggests that we might need to scale the sensor time series with respect to start of every individual engines time series.
- Scaling with respect to individual engines starting values allows us to bring all the engines time series to the same scale.