

fd001-baseline

June 14, 2024

1 Predictive Maintenance of Turbofan Jet Engine: Baseline Model

For the baseline of the model training, we will be using the most basic model without any pre-processing so that we can see how significant the improvement is after applying adequate data preprocessing and features engineering. - Baseline model: Linear Regression

2 Regression - RUL Prediction

2.1 1. Load Dataset

```
[1]: import warnings
warnings.filterwarnings("ignore")

import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns

from utils import read_dataset
```

```
[2]: train, test, test_rul = read_dataset("FD001")
train.shape, test.shape, test_rul.shape
```

```
[2]: ((20631, 26), (13096, 26), (100,))
```

```
[3]: train.head()
```

```
[3]:
```

	unit	time_cycles	op_setting_1	op_setting_2	op_setting_3	sensor_1	\
0	1	1	-0.0007	-0.0004	100.0	518.67	
1	1	2	0.0019	-0.0003	100.0	518.67	
2	1	3	-0.0043	0.0003	100.0	518.67	
3	1	4	0.0007	0.0000	100.0	518.67	
4	1	5	-0.0019	-0.0002	100.0	518.67	

	sensor_2	sensor_3	sensor_4	sensor_5	...	sensor_12	sensor_13	\
0	641.82	1589.70	1400.60	14.62	...	521.66	2388.02	
1	642.15	1591.82	1403.14	14.62	...	522.28	2388.07	
2	642.35	1587.99	1404.20	14.62	...	522.42	2388.03	

3	642.35	1582.79	1401.87	14.62	...	522.86	2388.08
4	642.37	1582.85	1406.22	14.62	...	522.19	2388.04

	sensor_14	sensor_15	sensor_16	sensor_17	sensor_18	sensor_19	\
0	8138.62	8.4195	0.03	392	2388	100.0	
1	8131.49	8.4318	0.03	392	2388	100.0	
2	8133.23	8.4178	0.03	390	2388	100.0	
3	8133.83	8.3682	0.03	392	2388	100.0	
4	8133.80	8.4294	0.03	393	2388	100.0	

	sensor_20	sensor_21
0	39.06	23.4190
1	39.00	23.4236
2	38.95	23.3442
3	38.88	23.3739
4	38.90	23.4044

[5 rows x 26 columns]

2.2 2. Establish Evaluation Metrics

Since we are solving a regression problem, we have choose two evaluation metrics to evaluate the trained models for estimating RUL. 1. **RMSE** - Root Mean Squared Error - one of the standard metrics for regression, it's a squared root of averaged squared difference between actual and predicted values. An important characteristic of RMSE is that it penalizes larger errors more.

$$RMSE = \sqrt{\frac{1}{n} \sum_{i=1}^n (y_i - \hat{y}_i)^2}$$

2. **MAE** - Mean Absolute Error - an average of absolute difference between actual and predicted values. MAE uses the same scale as the data and it's more robust to outliers.

$$MAE = \frac{1}{n} \sum_{i=1}^n |y_i - \hat{y}_i|$$

```
[4]: from sklearn.metrics import mean_squared_error, mean_absolute_error

def rul_evaluation_score(model, X, true_rul, metrics='all'):
    """
    Calculate evaluation metrics:
        1. rmse - Root Mean Squared Error
        2. mae - Mean Absolute Error

    Returns
    -----
    dict with metrics
```

```

'''
scores_f = {
    'rmse': lambda y_true, y_pred: np.sqrt(mean_squared_error(y_true,
↪y_pred)),
    'mae': mean_absolute_error
}

pred_rul = model.predict(X)

def calculate_scores(metrics_list):
    return {m: scores_f[m](true_rul, pred_rul) for m in metrics_list}

if metrics == 'all':
    return calculate_scores(scores_f.keys())
elif isinstance(metrics, list):
    return calculate_scores(metrics)

```

2.3 3. Establish Evaluation Methodology

The most important aspect to discuss in the cross-validation is that the the same engine cannot appear in 2 different folds. In time-series data or data where observations are grouped, standard k-fold cross-validation can lead to **data leakage**. This happens because the same engine can appear in both training and validation sets, which violates the assumption that the training and validation sets should be independent.

To avoid this, **GroupKFold** cross-validation method will be used, which ensures that the same group (engine unit) does not appear in both training and validation sets.

```

[5]: from sklearn.model_selection import GroupKFold
     from sklearn.model_selection import cross_validate

```

```

[6]: class CustomGroupKFold(GroupKFold):
     '''
     CV Splitter which drops validation records with
     RUL values outside of test set RULs ranges
     '''
     def split(self, X, y, groups):
         splits = super().split(X, y, groups)

         for train_ind, val_ind in splits:
             yield train_ind, val_ind[(y[val_ind] > 6) & (y[val_ind] < 135)]

```

```

[7]: def evaluate(model, X, y, groups, cv,
     scoring=['neg_root_mean_squared_error', 'neg_mean_absolute_error'],
     n_jobs=None,
     verbose=False):
     '''

```

```

Evaluate a model with Cross-Validation
'''
cv_results = cross_validate(
    model,
    X=X,
    y=y,
    groups=groups,
    scoring=scoring,
    cv=cv,
    return_train_score=True,
    return_estimator=True,
    n_jobs=n_jobs,
    verbose=verbose
)

for k, v in cv_results.items():
    if k.startswith('train_') or k.startswith('test_'):
        k_sp = k.split('_')
        print(f'[{k_sp[0]}] :: {" ".join(k_sp[2:])} : {np.abs(v.mean()):.
↪2f} +- {v.std():.2f}')
return cv_results

```

2.4 4. Build Baseline Model and Cross-Validate

```

[8]: from sklearn.pipeline import Pipeline
from sklearn.preprocessing import StandardScaler
from sklearn.linear_model import LinearRegression
from sklearn.compose import make_column_selector

from utils import calculate_RUL

```

```

[9]: get_ftr_names = make_column_selector(pattern='sensor')

```

```

[10]: baseline_model = Pipeline([
    ('scaler', StandardScaler()),
    ('model', LinearRegression())
])

cv_result = evaluate(
    baseline_model,
    X=train[get_ftr_names(train)].values,
    y=calculate_RUL(train, upper_threshold=135),
    groups=train['unit'],
    cv=CustomGroupKFold(n_splits=5)
)

```

```

[test] :: root mean squared error : 23.42 +- 1.54

```

```
[train] :: root mean squared error : 23.80 +- 0.35
[test] :: mean absolute error : 19.32 +- 1.25
[train] :: mean absolute error : 19.42 +- 0.35
```

2.5 5. Evaluate Baseline Model on Test Set

```
[11]: # Train model on the whole dataset
baseline_model.fit(
    X=train[get_ftr_names(train)].values,
    y=calculate_RUL(train, upper_threshold=135)
)

[11]: Pipeline(steps=[('scaler', StandardScaler()), ('model', LinearRegression())])

[13]: # Choose only the last cycle for each unit in the test set
X_test = test.groupby('unit').last().reset_index()


[14]: # Evaluate on the test set
rul_evaluation_score(baseline_model, X_test[get_ftr_names(test)], test_rul)

[14]: {'rmse': 22.368523506668193, 'mae': 17.88996899556987}
```

The results above will be used as our baseline when selecting the best ML models.

```
[15]: # plot the result
def plot_rul(y_true, y_pred):
    plt.figure(figsize=(12, 5))
    plt.plot(y_true, label='True RUL')
    plt.plot(y_pred, label='Predicted RUL')
    plt.legend()
    plt.show()

plot_rul(test_rul, baseline_model.predict(X_test[get_ftr_names(test)]))
```



The graph above will be compared to the selected model. We want to see whether the predicted RUL closely match with the true RUL or not.

[16]: *# Print the predicted value vs true value*

```
pd.DataFrame({
    'True RUL': test_rul,
    'Predicted RUL': baseline_model.predict(X_test[get_ftr_names(test)])
})
```

```
[16]:
```

	True RUL	Predicted RUL
0	112.0	127.233806
1	98.0	103.124768
2	69.0	75.339804
3	82.0	84.046613
4	91.0	93.794439
..
95	137.0	132.722916
96	82.0	90.107619
97	59.0	84.319839
98	117.0	129.667877
99	20.0	10.099036

[100 rows x 2 columns]